

# **Introduction to Machine Learning**

# Machine Learning is...

- Machine learning, a branch of artificial intelligence, concerns the **construction and study of systems that can learn from data**.
- The phrase “machine learning” describes a **wide diversity of algorithms and techniques**.

A big, expanding collection of algorithms and principles that analyze vast quantities of training data in order to extract meaning from it.

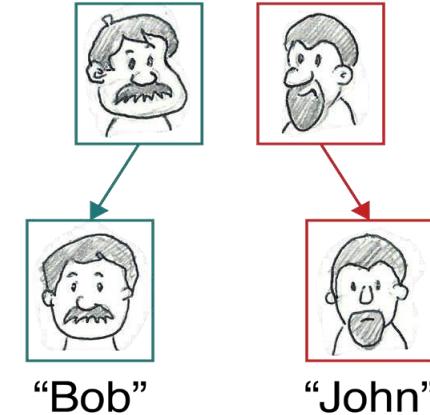
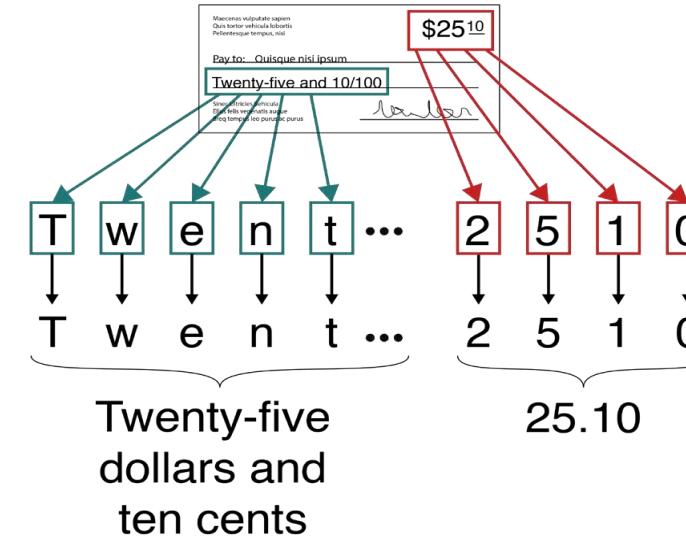
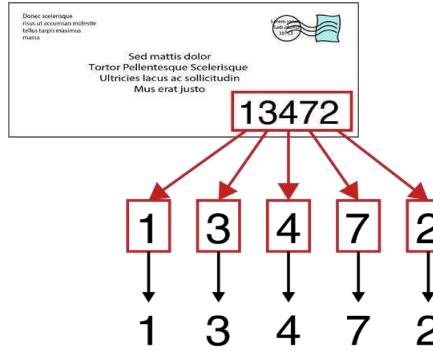
- The phrase machine learning describes a growing body of techniques that all have one goal: discover **meaningful information from data**.



# Data

- Data can be
  - **raw numbers** (like stock prices on successive days, the mass of different planets, the heights of people visiting a county fair).
  - **sounds** (the words someone speaks into their cell phone),
  - **pictures** (photographs of flowers or cats),
  - **words** (the text of a newspaper article or a novel),
  - **or anything else that we want to investigate.**
- “**Meaningful information**” is whatever we can extract from the data that will be useful to us in some way.
- We decide what’s meaningful to us, and then we design an algorithm to find as much of it as possible from our data.

# Example applications that use machine learning to extract meaning from data



- Left: Getting a zip code from an envelope.
- Middle: Reading numbers & letters on a cheque.
- Right: Recognizing faces from photos.

# Expert systems can also find meaning from data

- Expert systems was an early approach to finding the meaning that's hiding inside of data.
- Idea:
  - Study what human experts know and automate that.
  - Make a computer mimic the human experts it was based on.
- Create a rule-based system: a large number of rules for the computer to imitate human experts.
- Example: Recognize zip codes. 7's have a horizontal line at the top, and a diagonal line that starts at the right edge of the horizontal line and moves left and down. Some people put a bar through the middle of their 7's. So now we add another rule for that special case.

# But handcrafting rules is a tough job!

---

- This process of hand-crafting the rules to understand data is called **feature engineering**
- This term is also used to describe when we use the computer to find these features for us.
- It's easy to overlook one rule, or even lots of them. It's a tough job !

# How does ML compare with Expert systems?

---

- **Expert Systems:** Difficult to manually find right set of rules, & make sure they work properly across a wide variety of data. These difficulties have doomed expert systems.
- **ML Systems:** Beauty is they learn a dataset's characteristics automatically.
- Don't have to tell an algorithm how to recognize a cat or dog, because system figures that out for itself.
- **Flip side of ML:** To do its job well, ML system often needs a lot of data. *Enormous amounts of data.*

# Common threads in ML applications

---

**Sheer volume of work involved, and its painstaking detail.**

- We have millions of data to examine, and we want to extract some meaning from every one of them.
- Why can't humans do it ?
  - Humans get tired, bored, and distracted,
- What about Computers?
  - computers just plow on steadily and reliably.

**ML has ability to extract meaningful information quickly, so are used in many fields.**

# Why recent explosion in ML ?

---

- Why has machine learning has exploded in popularity and applications in the last few years?
- Couple of big reasons:

A. **Flood of data:** provided by the Internet has let these tools extract a lot of meaning from a lot of data.

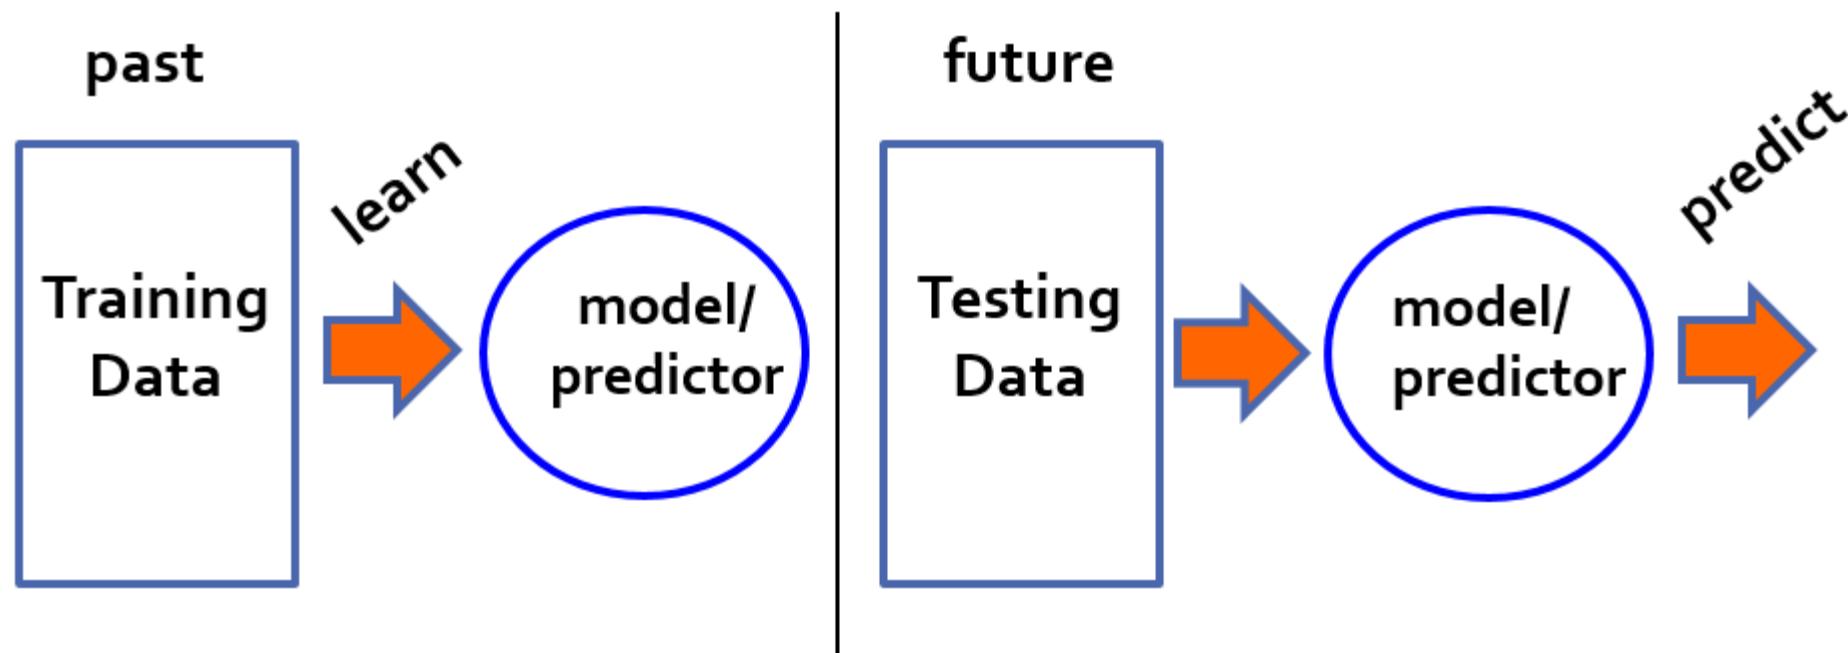
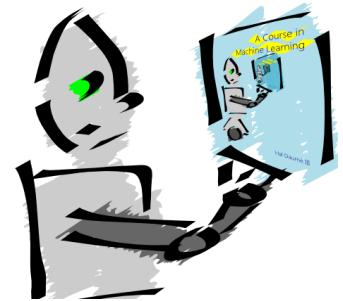
Example: Online companies make use of every interaction with every customer to accumulate more data. Then they use it as input to ML algorithms, getting more information about customers.

B. **Increased Computing power - GPUs**

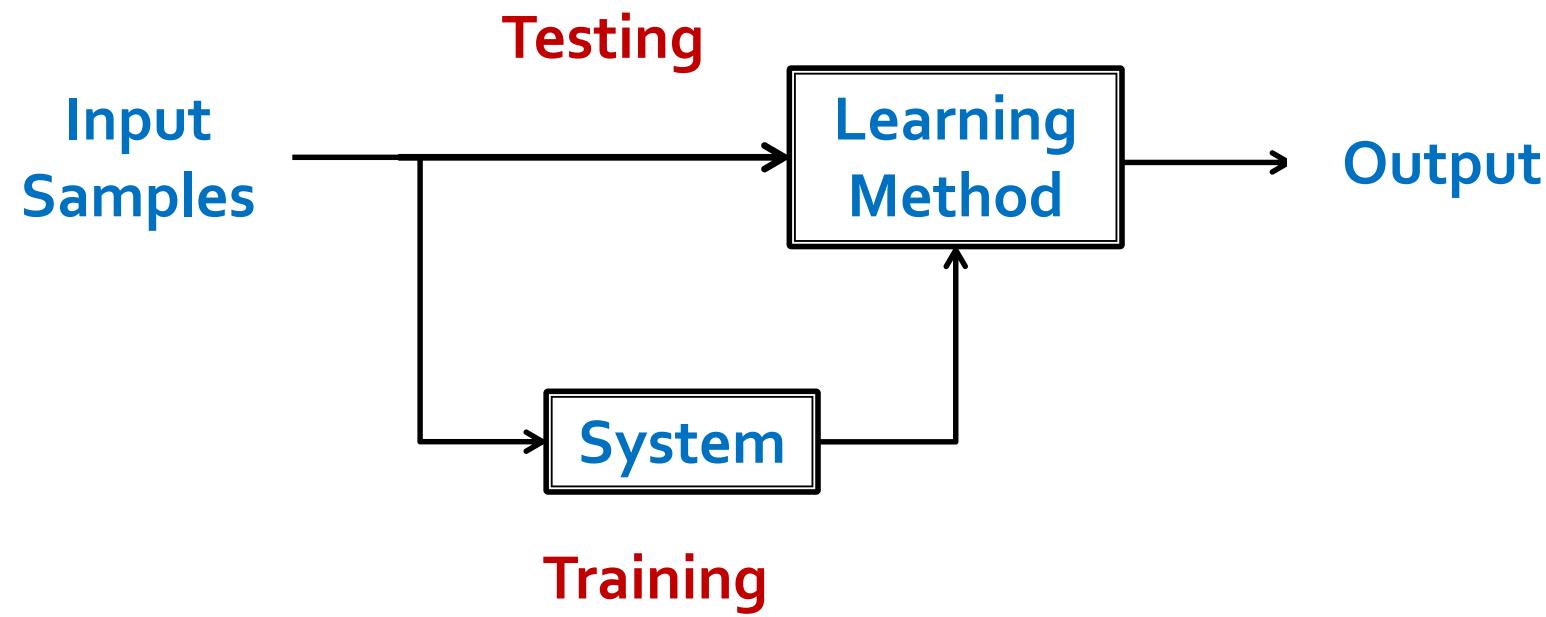
# Machine Learning is...

Machine learning is about predicting the future based on the past.

-- Hal Daume III



# Learning system model



# Performance

---

- There are several factors affecting the performance:
  - Types of **training** provided
  - The form and extent of any initial **background knowledge**
  - The type of **feedback** provided
  - The **learning algorithms** used

# What's a classifier ?

---

- A classifier assigns a label to each sample describing which category or class, that sample belongs to.
- Example of classifiers
  - If the input is a song, classifier assigns the label as the genre (e.g., rock or classical).
  - If it's a photo of an animal, the classifier assigns the label as the name of the animal shown (e.g., a tiger or an elephant).
  - In mountain weather for hiking, classifier may label the hiking experience into 3 categories: Lousy, Good, and Great.

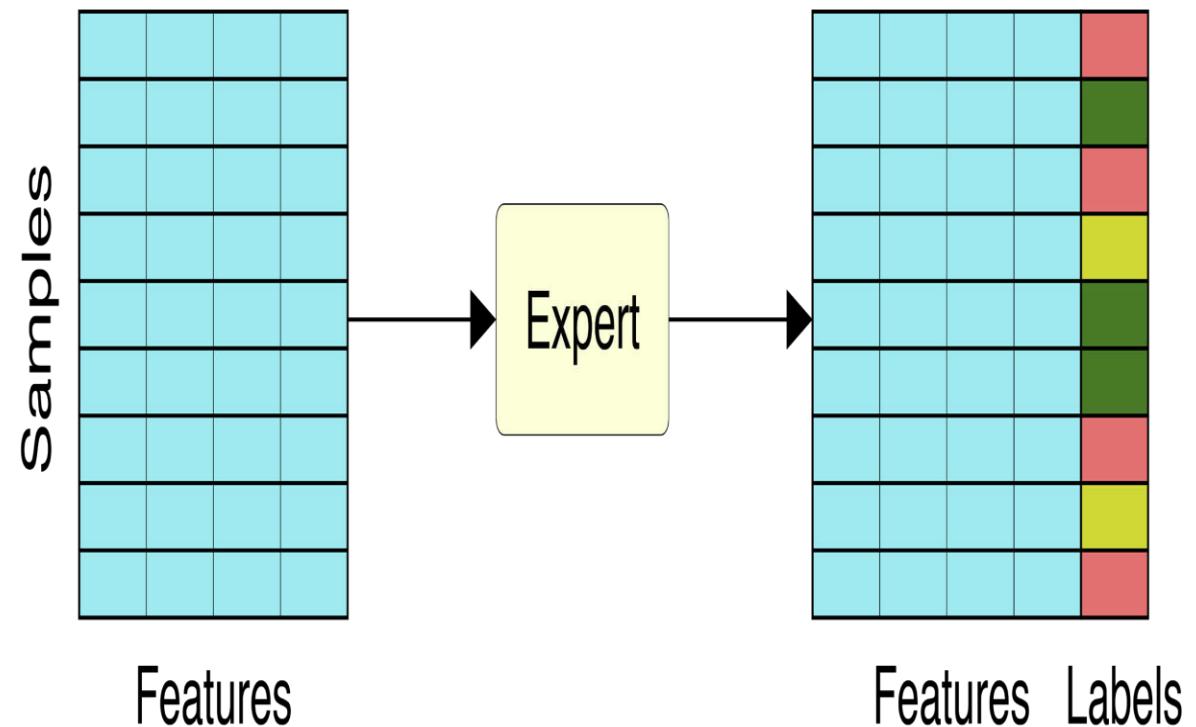
# Example of Samples, Features, Labels

---

- Weather measurements on a mountain for hiking
- Sample is weather at a given moment.
- Features are measurements: temperature, wind speed, humidity, etc.
- Hand over each sample (with a value for each feature) to a human expert.
- Expert examines features and provides a label for that sample.
- Expert's opinion, using a score from 0 to 100, tells how the day's weather would be for good hiking.
- Labels can be "Lousy", "Good", "Excellent" (weather for hiking)

# Example of Samples, Features, Labels (Contd.)

- To label a dataset, we start with a list of samples, or data items.
- Each sample is made up of a list of features that describe it.
- We give the dataset to a human expert, who examines the features of each sample one by one, and assigns a label for that sample.



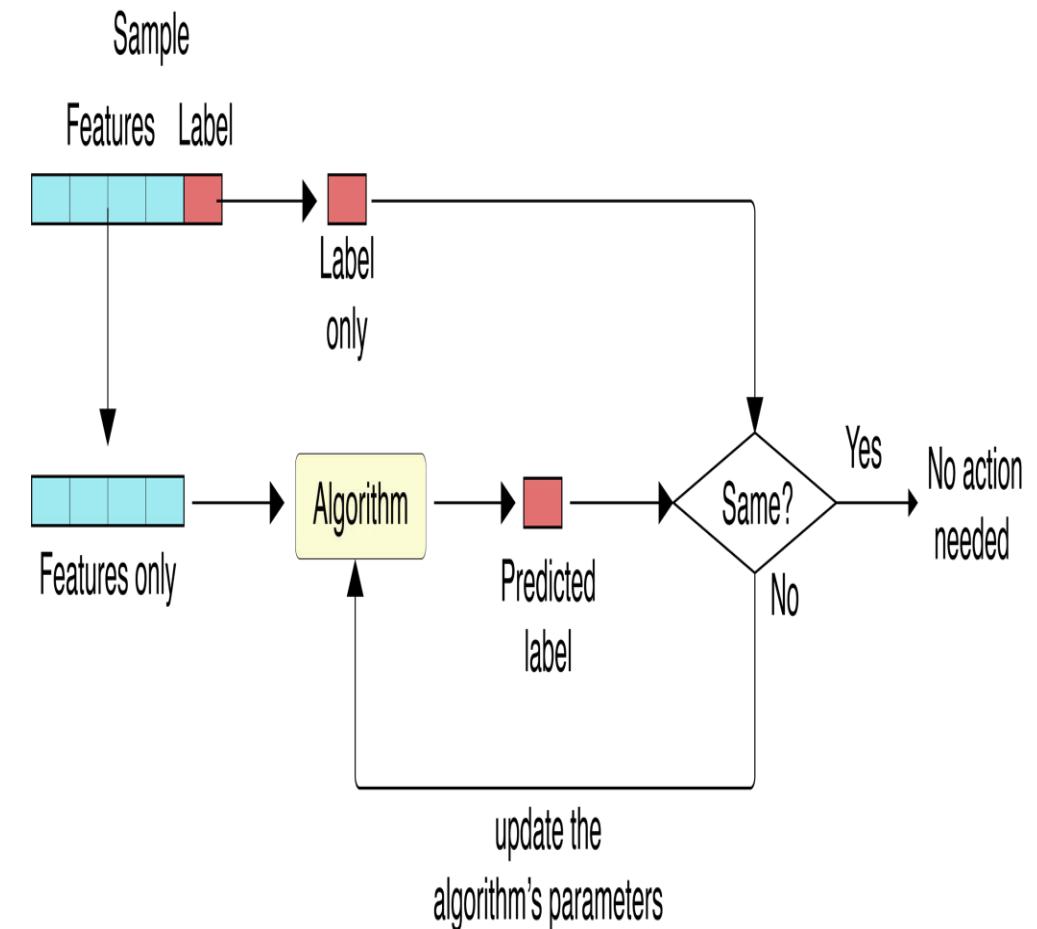
# A Computerized Learning Strategy

---

- First, collect as much data as possible.
- Call each piece of **observed data** (say, the weather at a given moment) as **sample**.
- Call the **names of the measurements** that make it up (the temperature, wind speed, humidity, etc.) as **features**.
- Hand over each sample (with a value for each feature) to a human expert.
- Expert examines features and provides a **label** for that sample.
- Example: if our sample is a **photo**, the label might be the **name of the person** or the type of **animal in the photo**.

# A Computerized Learning Strategy

- Figure shows the idea of a learning strategy - one step of training or learning.
- Split the sample's features and its label. From the features, algorithm predicts a label.
- Compare prediction with truth label.
- If predicted label matches truth label, don't do anything.
- Otherwise, tell algorithm to update itself
- The process is basically trial and error.

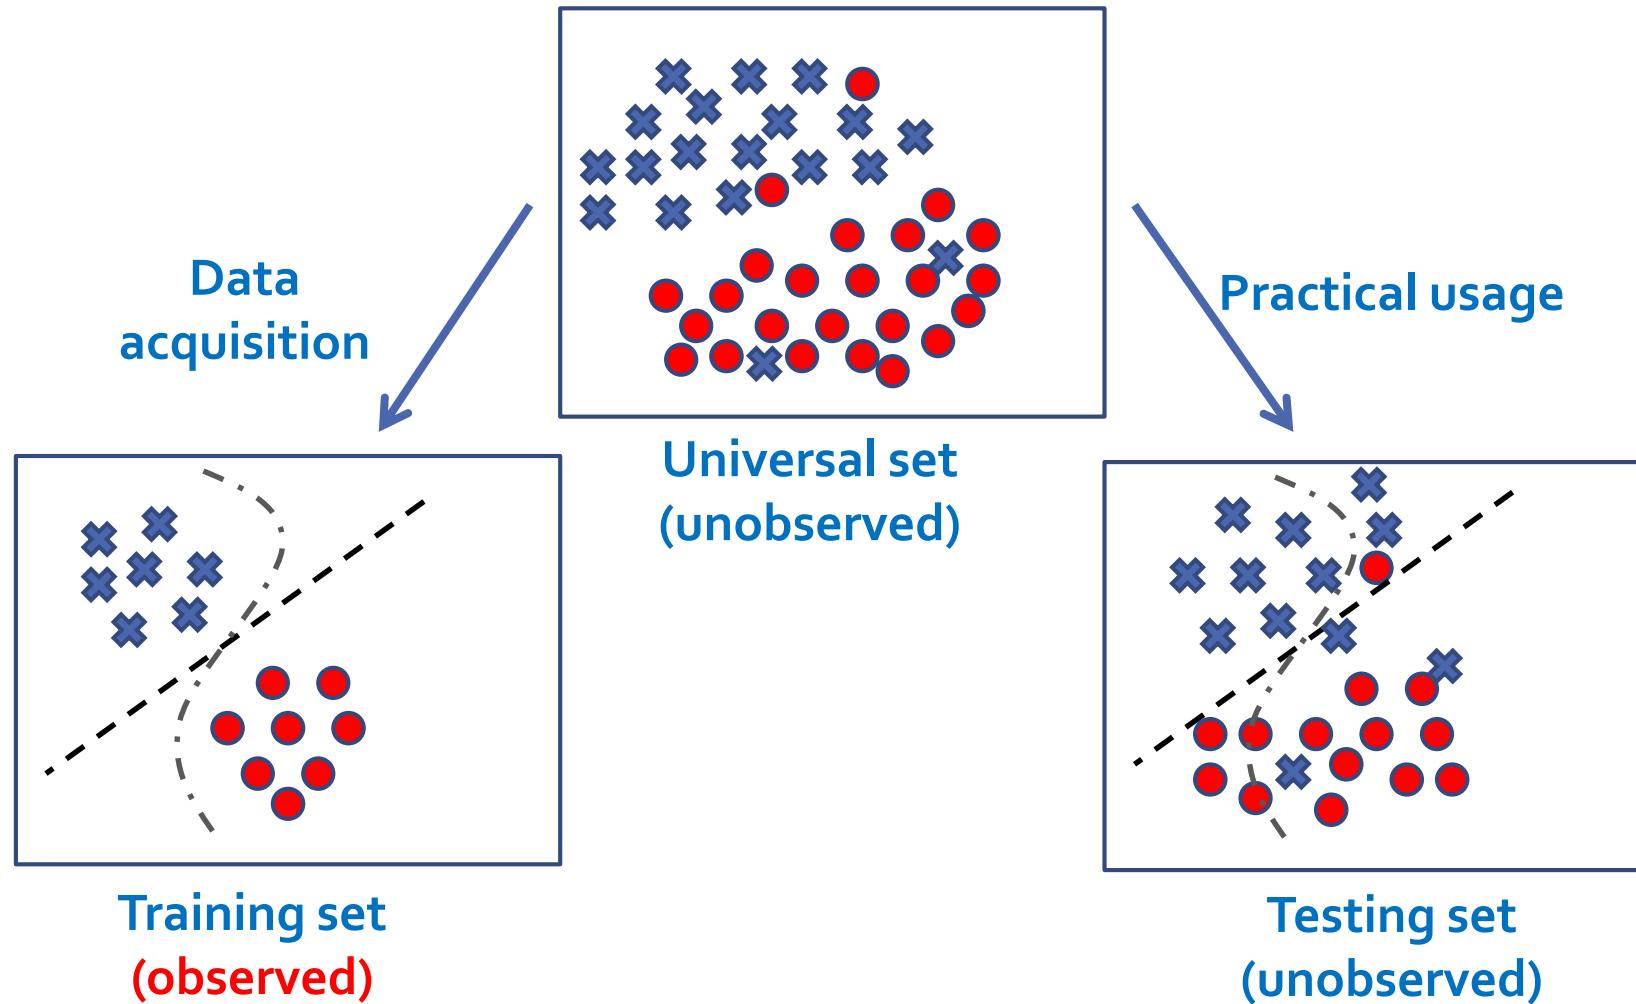


# First, split Data for Training & Validation

---

- First, set aside some of these labeled samples for time being (use them later for validation).
- Give remaining labeled data to our computer, and ask it to find a way to come up with the right label for each input.
- We do not tell it how to do this.
- Instead, we give labelled data to an algorithm with a large number of parameters it can adjust (perhaps even millions of them).
- Different types of learning will use different algorithms.

# Training and Testing



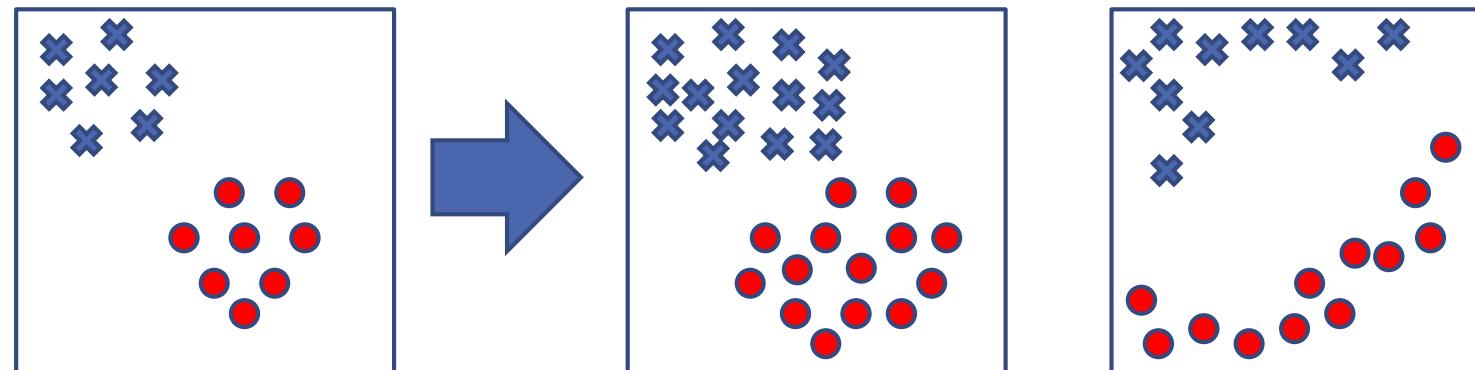
# Training data

---

- Training is the process of taking a system that's been initialized with default or random values, and reconfiguring it so that it's tuned to the data we want to understand.
- Training data is data used to train the system.
- Real-world data is data the classifier will see in the real world, once it's been released, or deployed.
- Real-world data will only arrive after system has been deployed.

# Training and testing

- **Training** is the process of making the system able to learn.
  - Training set and testing set come from the same distribution
  - Need to make some assumptions or bias



# Parameters and hyper parameters

---

- Learning algorithm modifies its own parameter values, over time.
- Learning algorithm are also controlled by values that we set (such as the learning rate we saw above).
- These are called **hyper parameters**.

## Difference between parameters and hyper parameters

- Computer adjusts its own parameter values during the learning process, while we specify the hyper parameters when we write and run our program.

# Training step and Learning rate

---

- Each algorithm **learns by changing the internal parameters** it uses to create its predictions.
- **Big change:** risk of changing them so much that it makes other predictions worse.
- Small change: Cause learning to run slower.
- We have to find by trial and error for each type of algorithm the **right trade-off between these extremes.**
- We call the amount of updating the **learning rate,**
- A **small learning rate is cautious and slow,**
- A **large learning rate speeds things up but could backfire.**

# Testing or Validation step

---

- We now return to the labeled data **kept aside** in the last section.
- This is called as **test data**.
- We evaluate how the system can **generalize** what it learned, by showing these samples that it's **never seen** before.
- This **test set** shows how the system performs on **new data**.

# Validation

---

- Q: What is validation?
  - A: Validation is **evaluating how well system is learning.**
- Q: How do you validate a system?
  - A: By showing it new data(from training set, not real-world)it hasn't seen before, and measuring how accurately it works
- Q: What happens after validation?
  - A: If predictions on validating data are good, deploy system.
    - - If not good, then go back and train some more.

# Cross-validation

---

- Cross-validation is a validation tool
- It squeezes every bit of learning from training data.
- Cross-validation is specially useful when dataset is small.

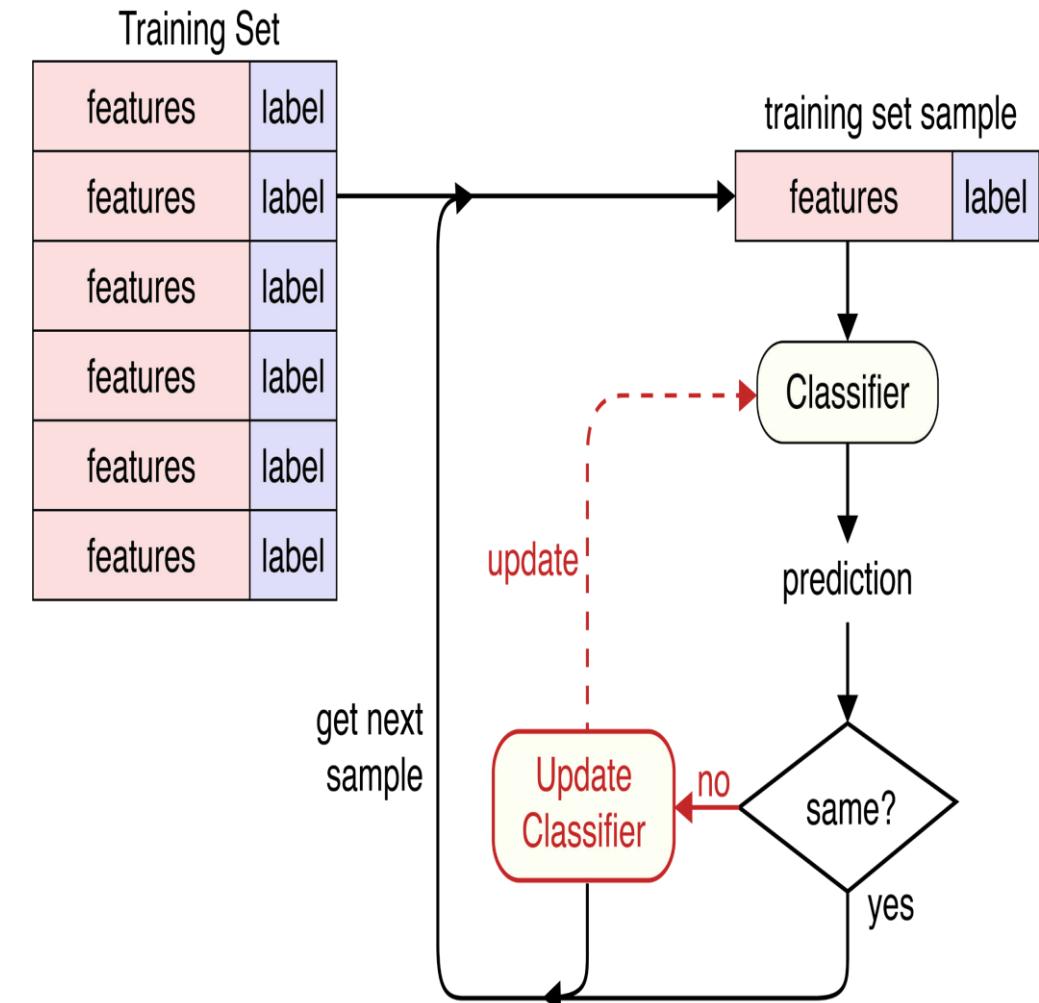
# Lets train and test a Categorizer (for example)

---

- We'll illustrate the ideas in this module using a supervised categorizer, which we'll teach with labeled data.
- In particular, we'll use a neural network type of classifier, which learns by repeated exposure to the training data.
- The techniques we'll discuss are general, and can be applied to all types of learners.

# Lets train and test a Categorizer (for example)

- Present each sample in training set to classifier, one at a time.
  - For each sample, give the sample's features and ask it to predict its category.
  - If prediction is correct, then move to next sample.
  - If prediction is wrong, feed classifier's output and correct label back to classifier.
  - ML algorithm uses correct label, predicted label, and the current state of classifier to modify the classifier's internal parameters - so that it's likely to predict correct label if it sees this sample again.



# How long to train? Epochs

---

- **Epoch** is a run through the entire training set.
- Usually you must train through many epochs.
- Keep training as long as the system is still learning and improving its performance on the test data.
- Stop **when system's performance is good enough for the task.**
- Stop **also if you are out of time.**

# Testing performance – separate issue performance on training data is not enough

- Aim: Find how well categorizer will perform on real-world data before it is deployed. Categorizer should meet a quality threshold.
- Q: How to estimate quality of predictions before categorizer is released?
- Q: Is it enough if our categorizer does really well on training data?
  - A: No. If we just use the quality of the classifier's results only from training data, we'll be misled.

# Testing performance – separate issue performance on training data is not enough

- Q: Why is good performance on training data alone not enough?
- A: Because we don't really know what the system learned from the training data.
- It might have learnt about some weird idiosyncrasies in training data and then used that as a rule, only to be foiled by new data that doesn't happen to have those quirks.

# Run experiments to check performance

---

- Q: How do we then evaluate performance of the categorizer?
- A: There's no way for us to know how good our categorizer is without trying it out and seeing.
- We must run experiments to see how well our system performs.
- The best way anyone has found to work out how well a system will do on new, unseen data is to give it new, unseen data and see how well it does.
- There's no shortcut to this kind of experimental verification.

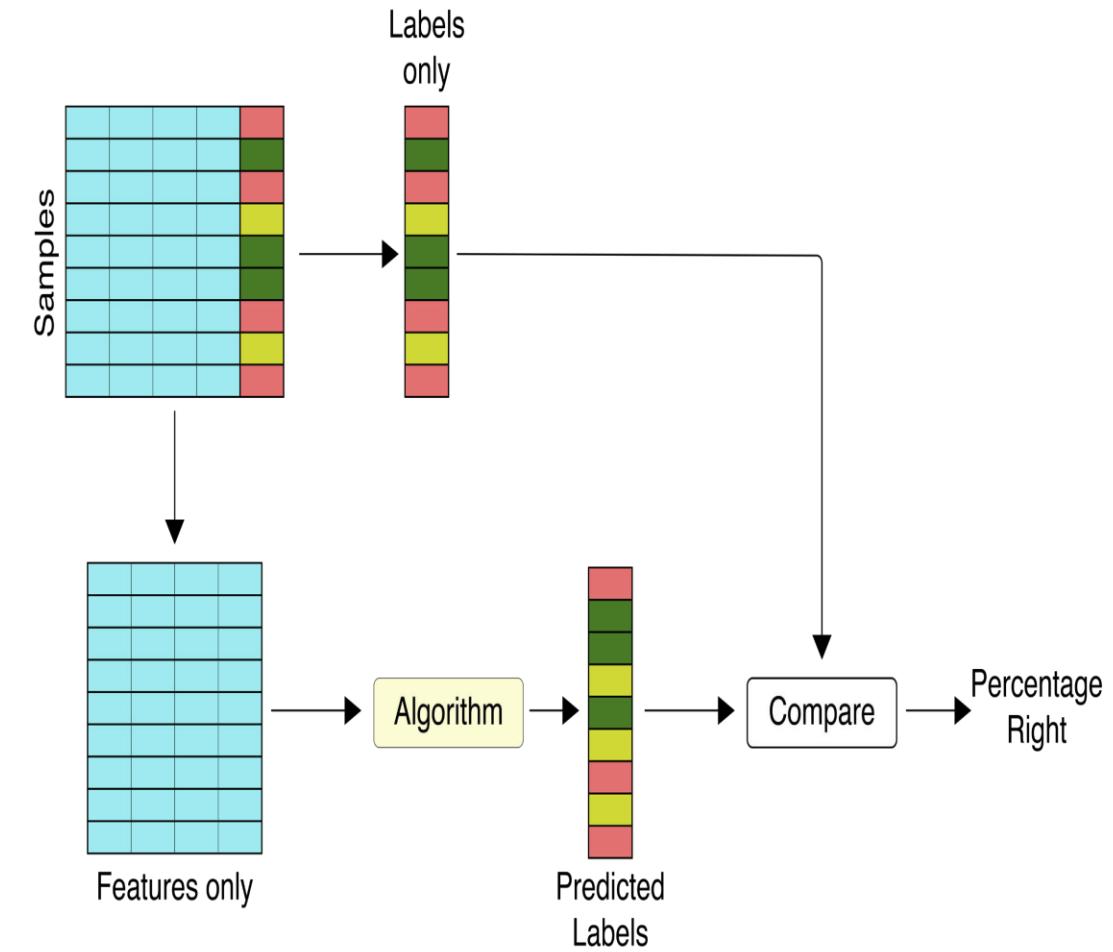
# Test data – a representation for real-world data

---

- **Test data is the new unseen data used for testing.**
- We hope test data represents real-world data that system will see when deployed.
- If performance on test data isn't good, then train it more.
- Irrespective of how well it might have done on the training data.

# Testing – Procedure for evaluating a classifier

- Split the test data (not training data) into features and labels.
- Algorithm predicts a label for each set of features.
- Compare predictions with the truth labels to get a measurement of accuracy.
- If it's good enough, deploy the system.
- If the results aren't good enough, go back and train some more.
- In this evaluation process there is no feedback and no learning.



# Never learn from test data

---

- Never learn from the test data.
- If categorizer learns from test data, it loses its special and valuable quality as a way to measure the performance of the system on new data.
- This would ruin the reliability of our measure of the general ability of our system.
- So, this would be “cheating,” because it’s already learned from test data.
- So keep test set hidden from the system until training is over, and then use it just once to estimate its performance.

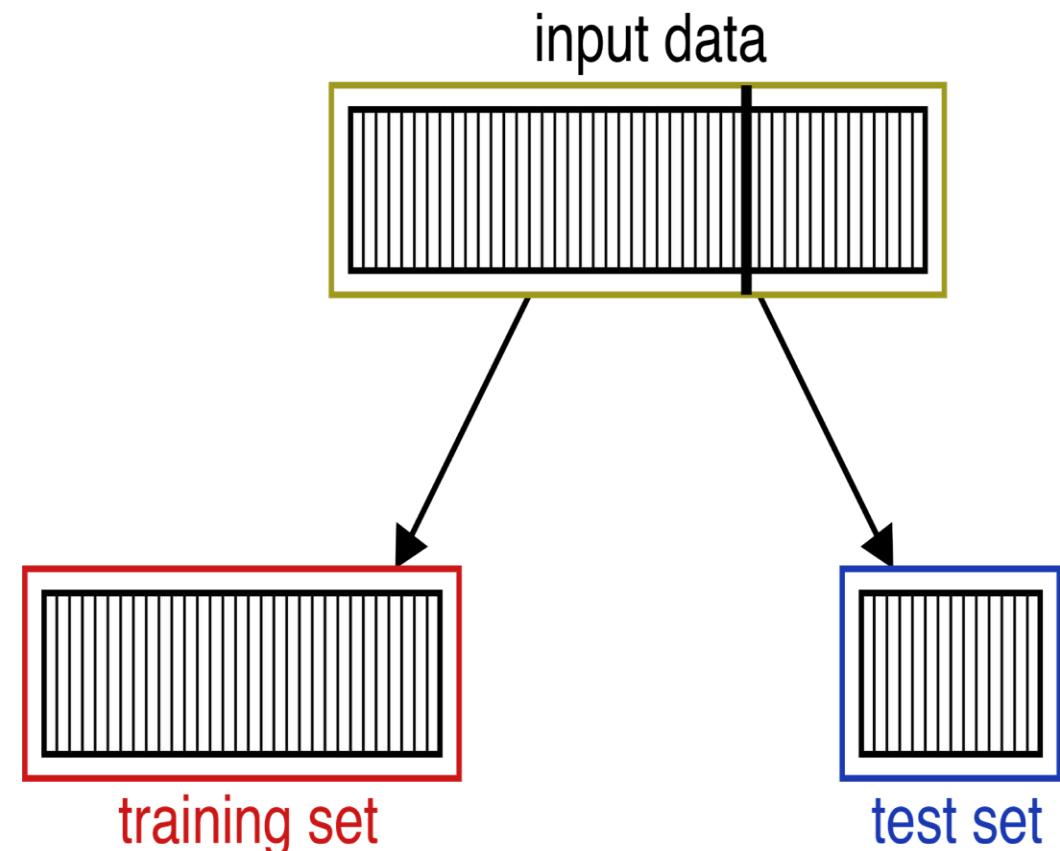
# Lookout for Data leaks

---

- Data leakage, or data contamination happens when system learns from test data.
- Watch out for this !
- Data leaks can sneak in, as they are hard to notice.

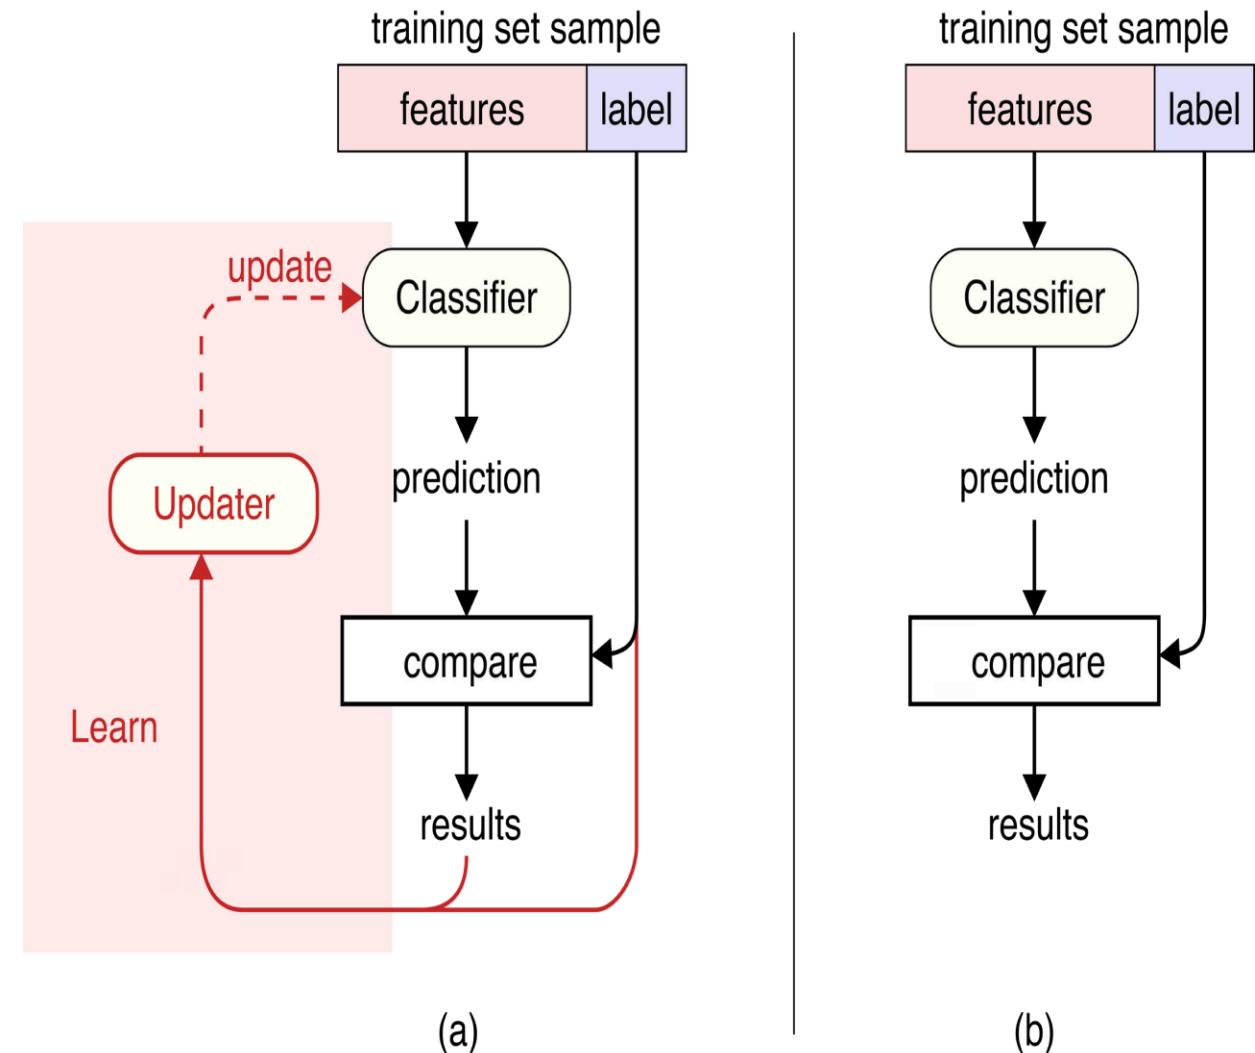
# How to create training and test data sets?

- Create test data by splitting original data collection into two pieces: training set and test set.
- Split to give about 75% of the samples to training set.



# Difference between training and testing

- In **training**, compare the prediction and the real label, and use any error between them to make classifier learn to do better.
- In **testing**, omit learning step.



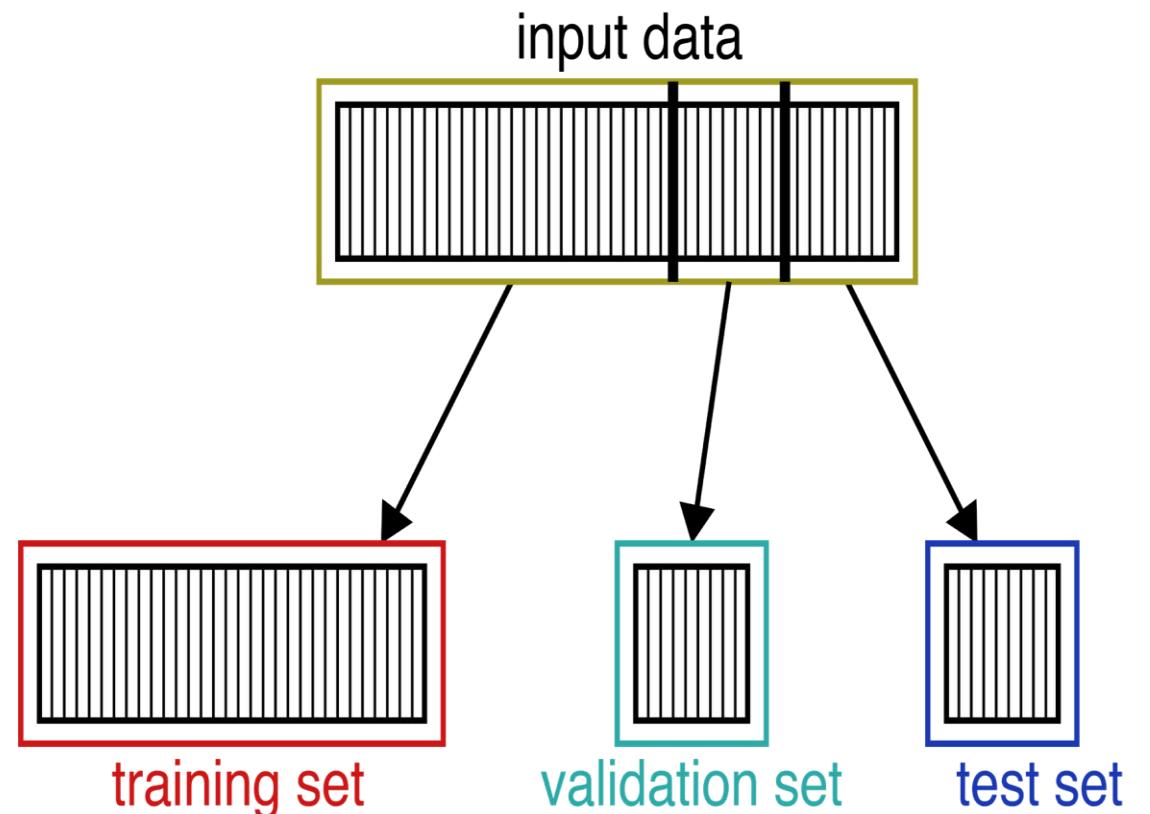
# To try hyperparameters, use validation data

---

- To try out many values of hyperparameters, create and use validation data.
- For each variation of hyperparameters, train on the training set, and evaluate the system's performance on the validation set.
- Select the hyperparameters that gave us the best results.
- Run that through the test set to evaluate the system's performance on data that it's never seen before.
- In other words, use the validation set as part of the learning process.
- The test set is still held aside for that one-time final evaluation.

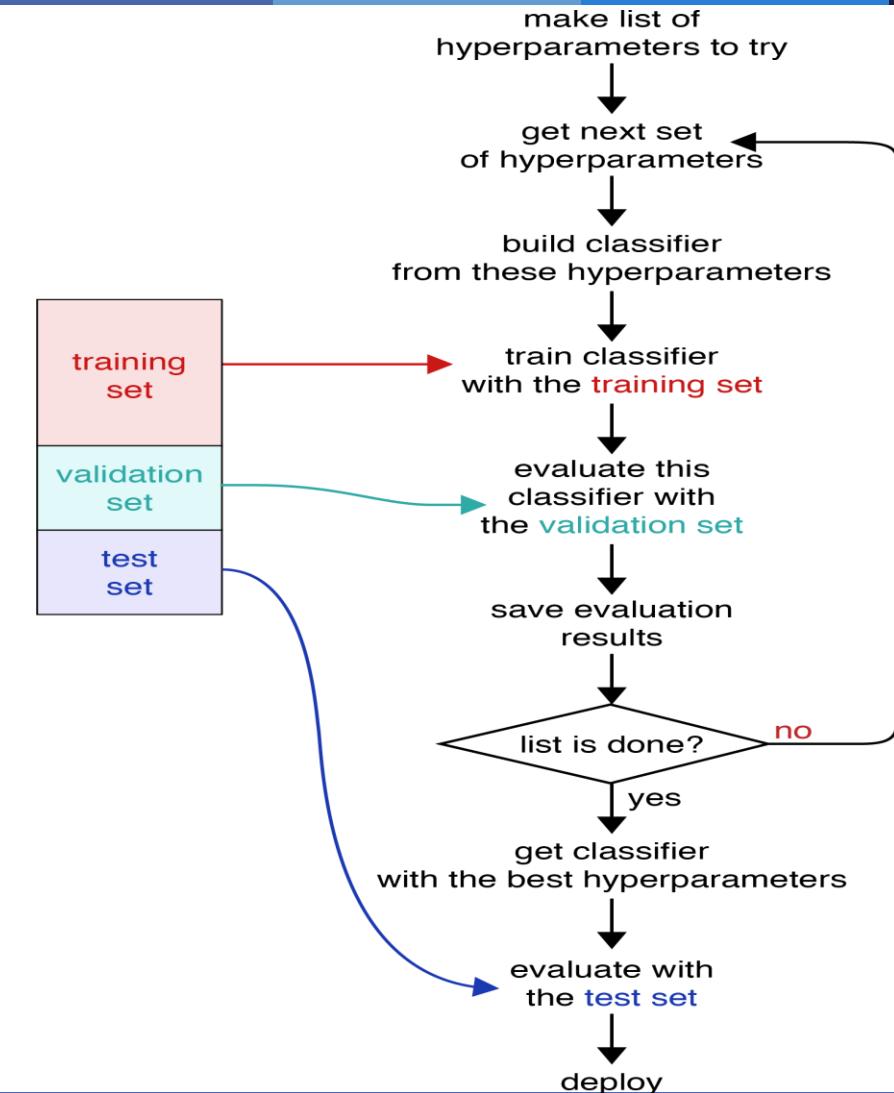
# Creating Validation data set

- Splitting input data into a training set, validation set, and test set.
- Split original data into three pieces, 60% for training, 20% for testing, and 20% for validation.



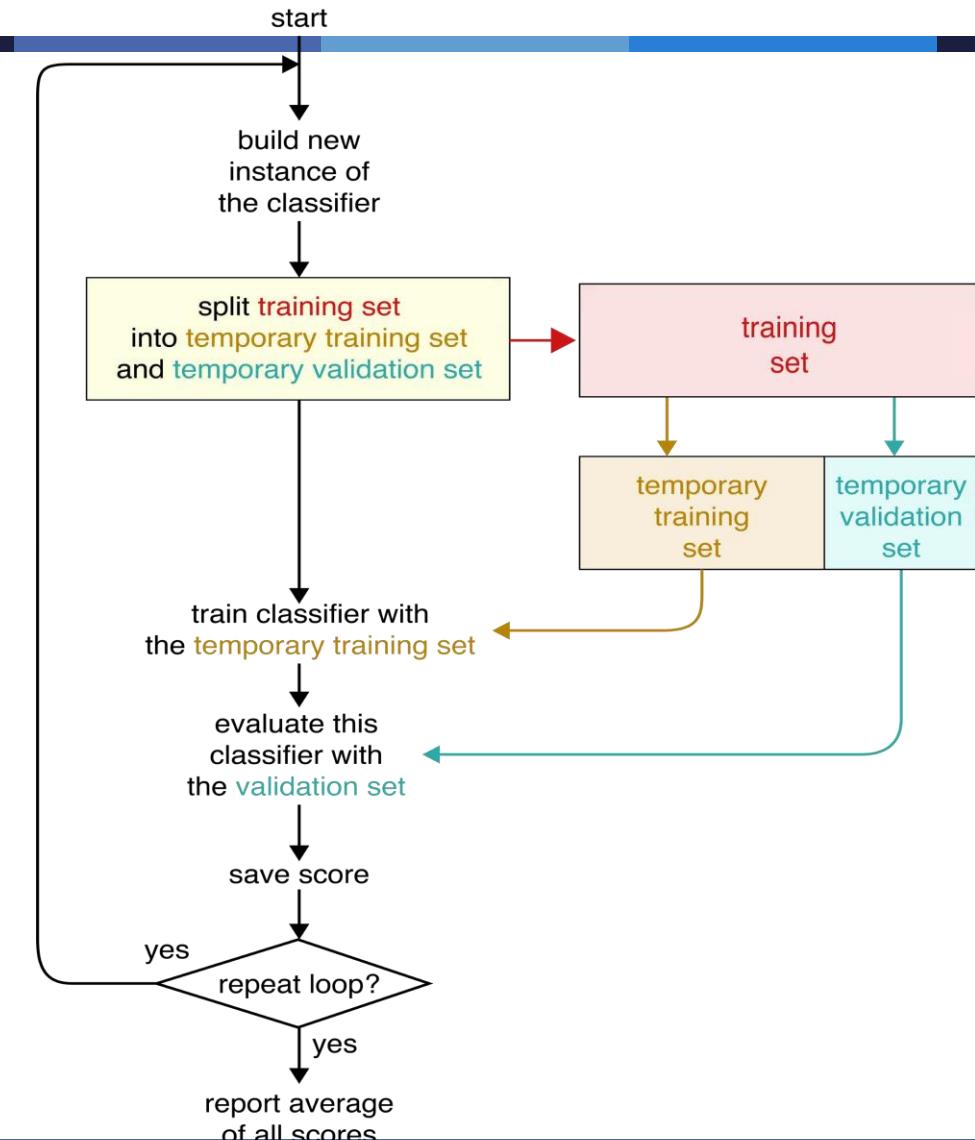
# Validation Process- for trying lots of hyperpairs

- Run the Loop:
  - Select a set of hyperparameters, Train the system,
  - Test it with validation set.
  - Repeat with new set of hyperpars
- Select set of hyperpars that gave best performance.
- Take that system, apply it to testing data. & check how good its predictions are.
- If good, deploy. Else, retrain
- We'll see next a more sophisticated method called cross-validation that does not re-use the same validation set over and over in this way.

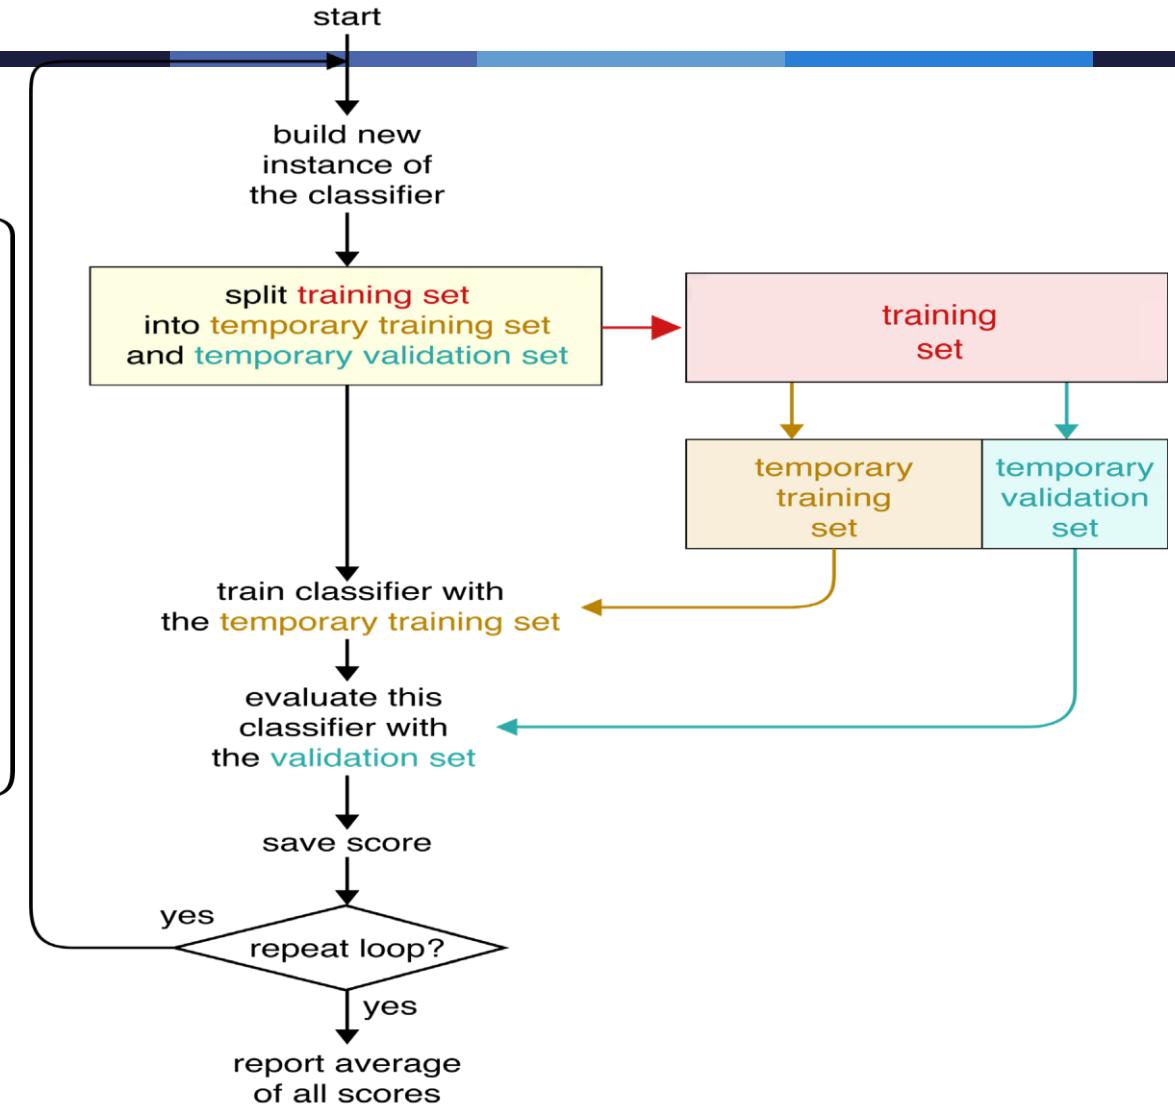
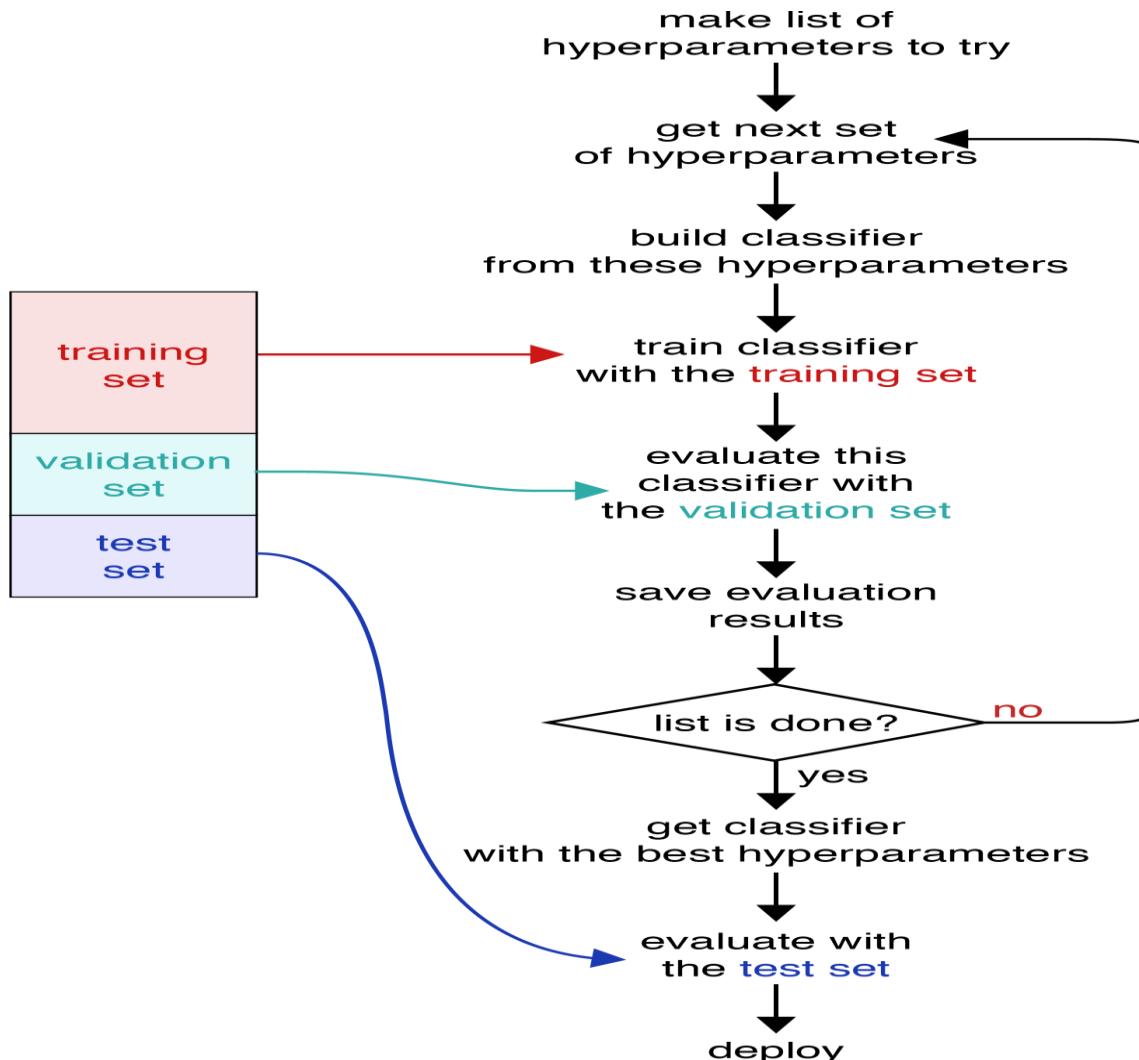


# Validation Process- for trying lots of hyperpairs

- Before starting, pull out the test data and set it aside.
- Work with rest of data- training data set
- Each time through loop, split training data into a temporary training set and temporary validation set.
- Begin by building a fresh instance of our classifier.
- Train on temporary training set, and evaluate with temporary validation set.
- Get score for current classifier's performance.
- Go through loop again, split data into new sets, different from any we've tried before.
- Loop over and over, splitting data into fresh sets, training and validating, and getting a score.
- Find average of all scores of our classifier.



# Inner loop of Validation- k-fold Cross validation - for hyperparms



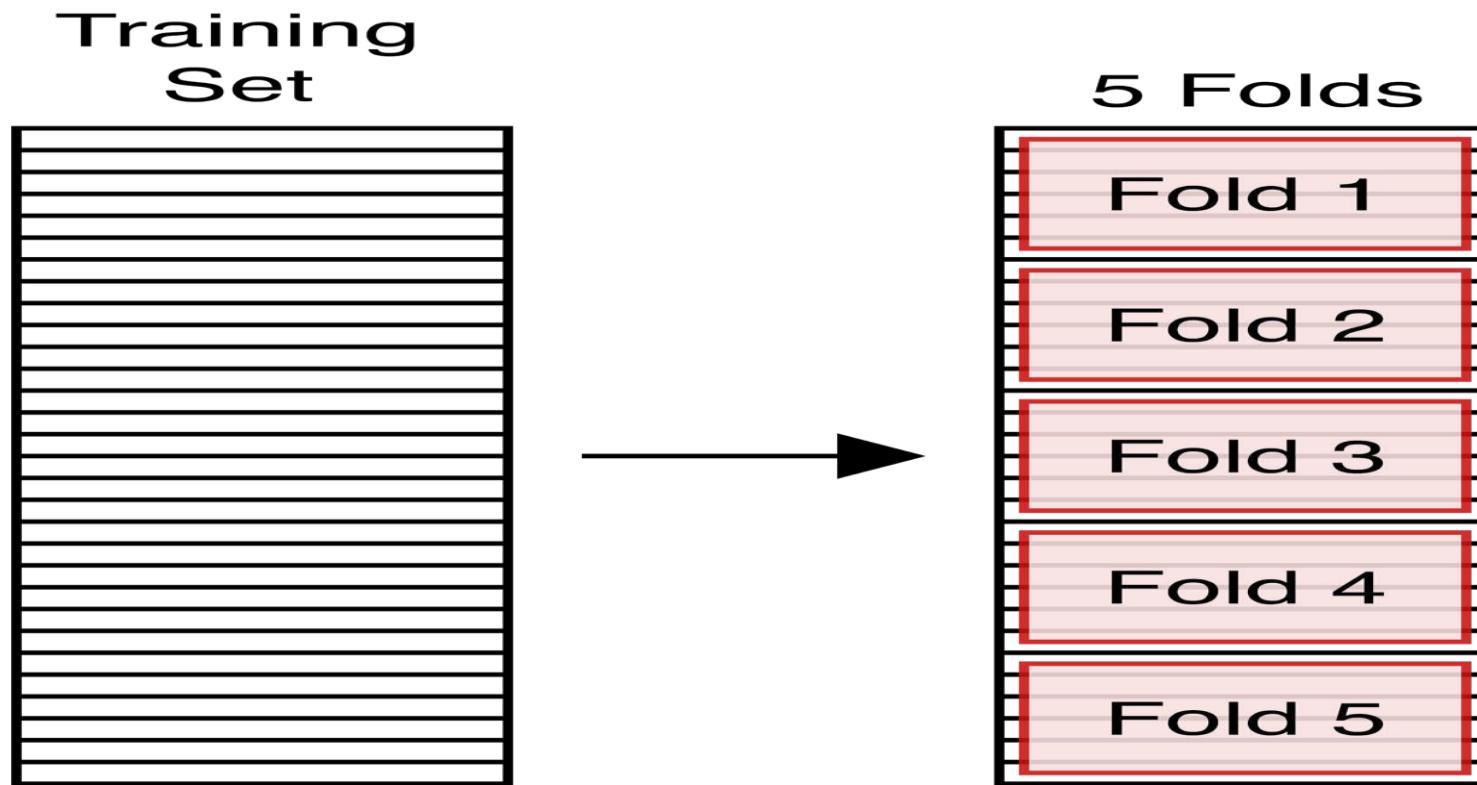
# K-fold cross-validation

---

- There are a variety of different algorithms available for constructing the temporary training and validation sets.
- Let's look at a popular approach for constructing temporary data sets. Called as **K-fold cross-validation**
- Popular way to build the temporary datasets for cross-validation is called k-fold cross-validation.
- For example, “2-fold cross-validation” or “5-fold cross-validation”. Typically, the value of k is the number of times we want to go through the loop.
- The algorithm starts before the cross-validation loop of Figure begins.
- We take our training data and split it up into a series of equal-sized chunks. We call each chunk a fold.

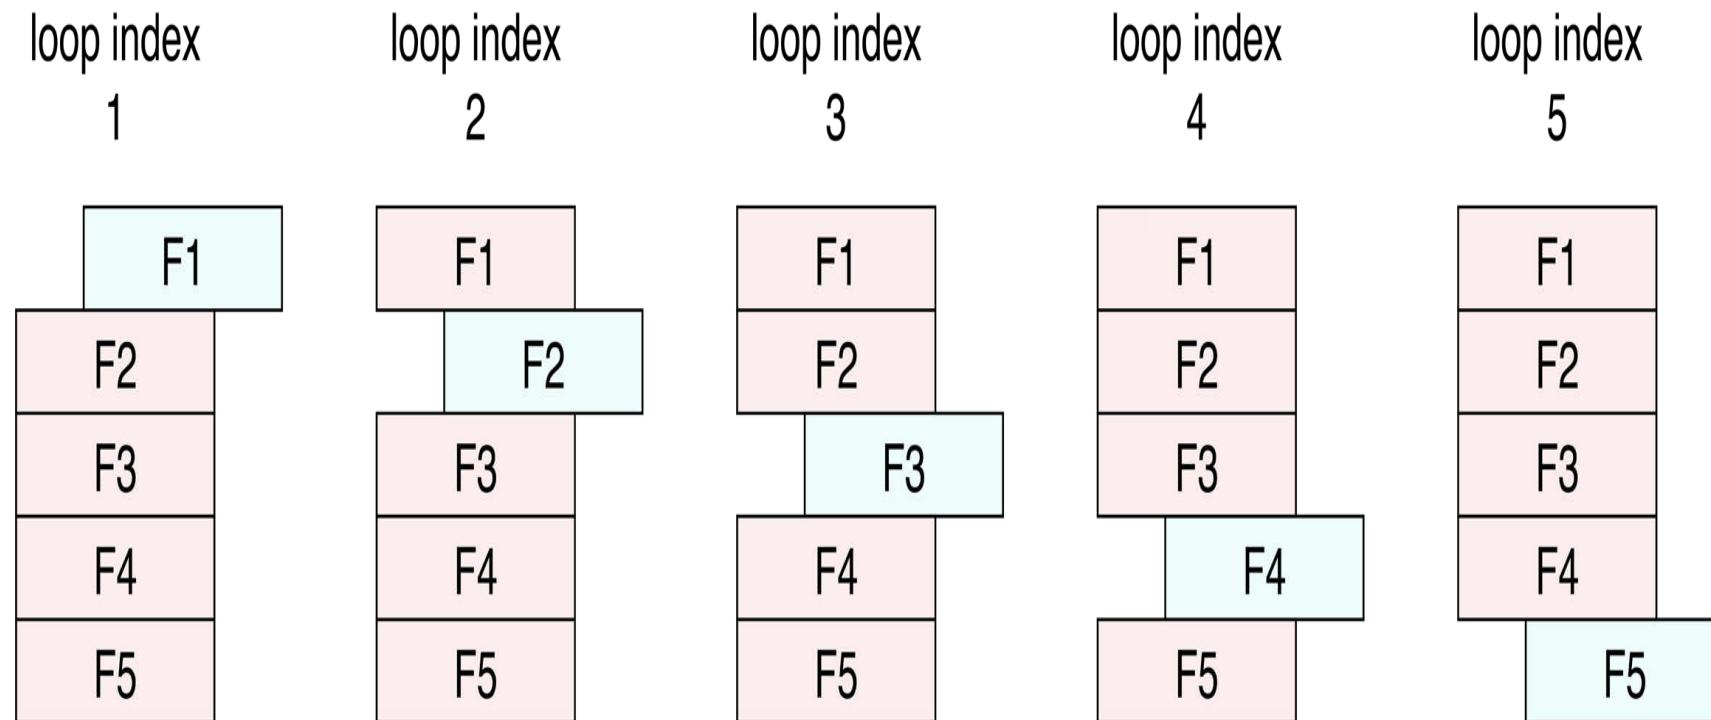
# 5-fold – splitting data set

Splitting our training set into 5 equally-sized folds, named Fold 1 through Fold 5.



# Example of 5-fold

- In each pass through the loop, we choose one fold for validation (in blue), and train with the others (in pink).
- If we go through the loop more than 5 times, we just repeat the pattern.



# How to Retrain, if testing results are not good

---

- Use again **original training set** data. Note that these are the same samples.
- **Shuffle** this data first - but no new information.
- Show every sample again, letting it learn along the way again.
- Computer learns over and over again from the very same data.
- Now, show test data set .
- Ask algorithm to predict labels for the test set again.
- If the performance isn't good enough, go back to original training set again, and then test again.
- Repeat this process often hundreds of times. Let it learn just a little more each time.
- Computer doesn't get bored or cranky seeing the same data over and over.

# Learning – Good and Bad News

---

- **Bad News:**
  - No guarantee that there's a **successful learning algorithm for every set of data**,
  - No guarantee that if there is one, we'll find it.
  - May not have enough **computational resources** to find the relationship between the samples and their labels.
- **Good news:**
  - Even without a mathematical guarantee, in practice we can often find solutions that generalize very well, sometimes doing even better than human experts.

# When do we deploy the System ?

---

- When the algorithm has learned enough to perform well enough on the test set that we're satisfied, we're ready to **deploy, or release**, our algorithm to the world.
- Users submit data and our system returns the label it predicts.
- That's how pictures of faces are turned into names, sounds are turned into words, and weather measurements are turned into forecasts.

# Machine Learning – Major categories-Algorithms

---

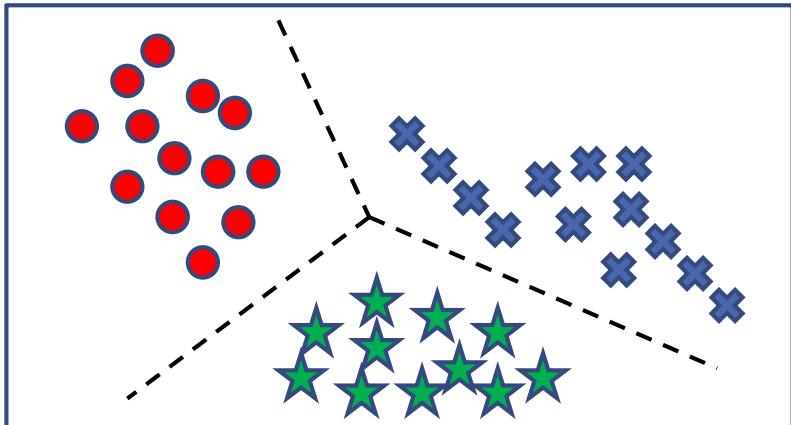
- The success of machine learning system also depends on the **algorithms**.
- The algorithms control the search to find and build the **knowledge structures**.
- The learning algorithms should extract **useful information** from training examples.

# Algorithms

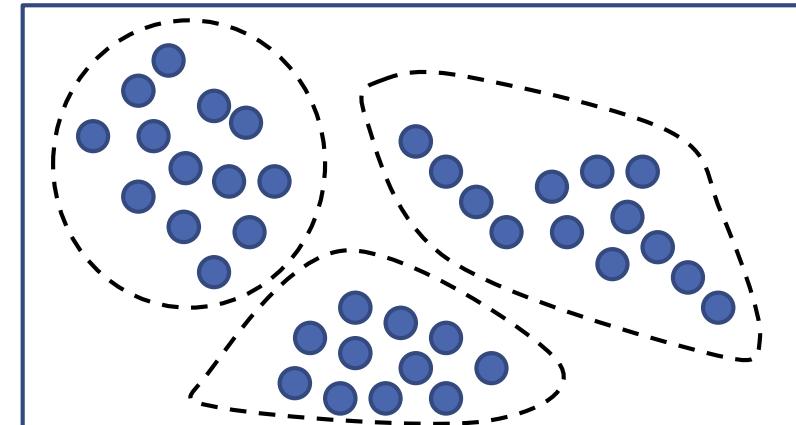
---

- **Supervised learning**
  - Prediction
  - Classification (discrete labels), Regression (real values)
- **Unsupervised learning**
  - Clustering
  - Probability distribution estimation
  - Finding association (in features)
  - Dimension reduction
- **Semi-supervised learning**
- **Reinforcement learning**
  - Decision making (robot, chess machine)

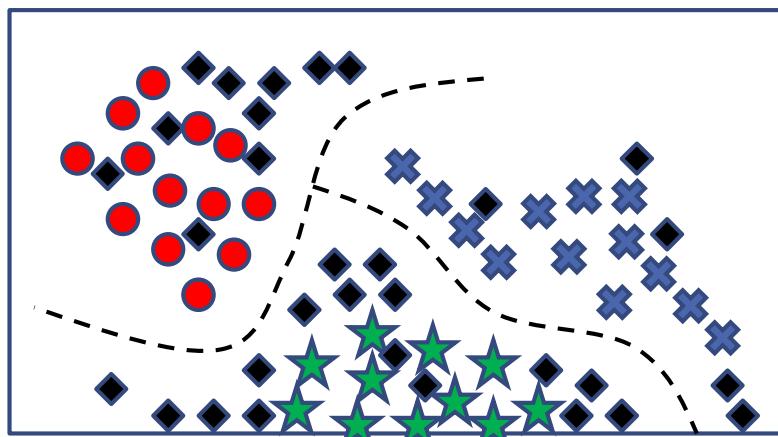
# Algorithms



Supervised learning



Unsupervised learning



Semi-supervised learning

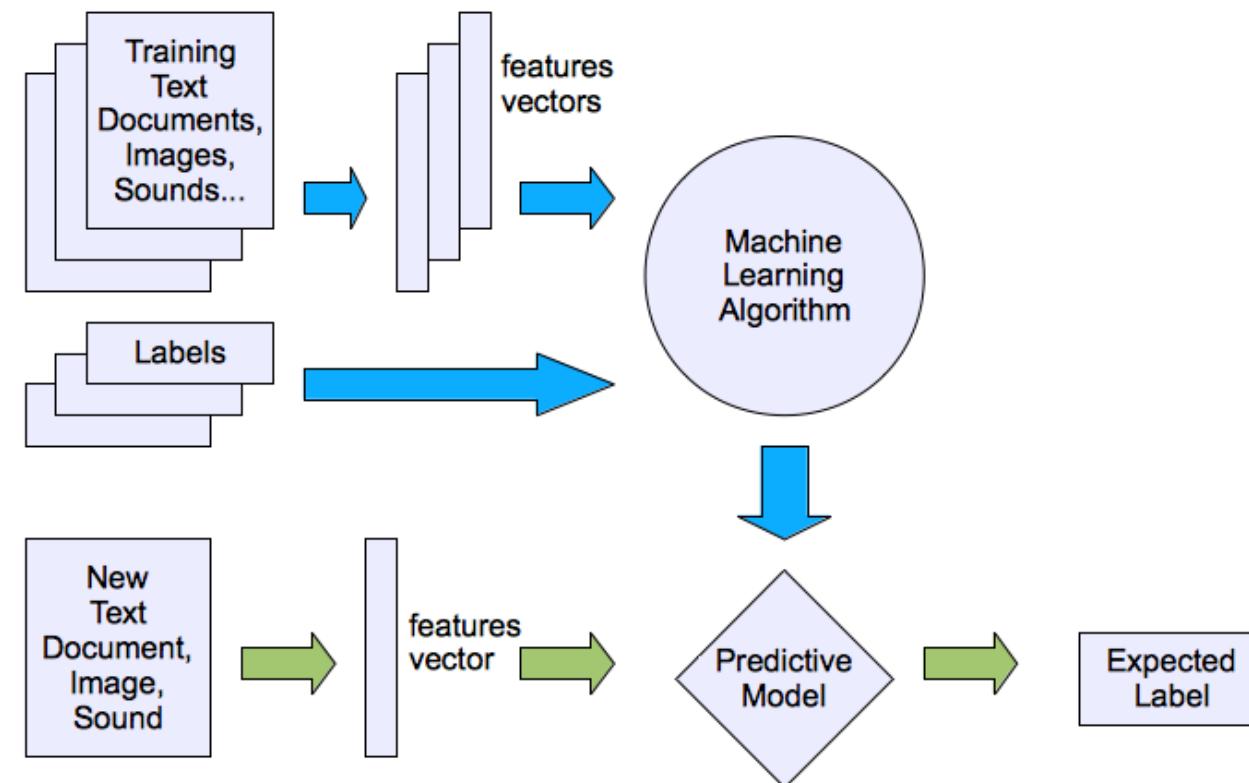
# Supervised Learning (ML)

---

- Supervised learning (SL) is done for samples with pre-assigned labels.
- Supervision comes from the labels.
- Labels guide the comparison step.
- There are two general types of supervised learning, called classification and regression.
- Classification: look through a given collection of categories to find the one that best describes a particular input.
- Regression: take a set of measurements and predict some other value

# Machine learning structure

- Supervised learning



# SL – Classification

---

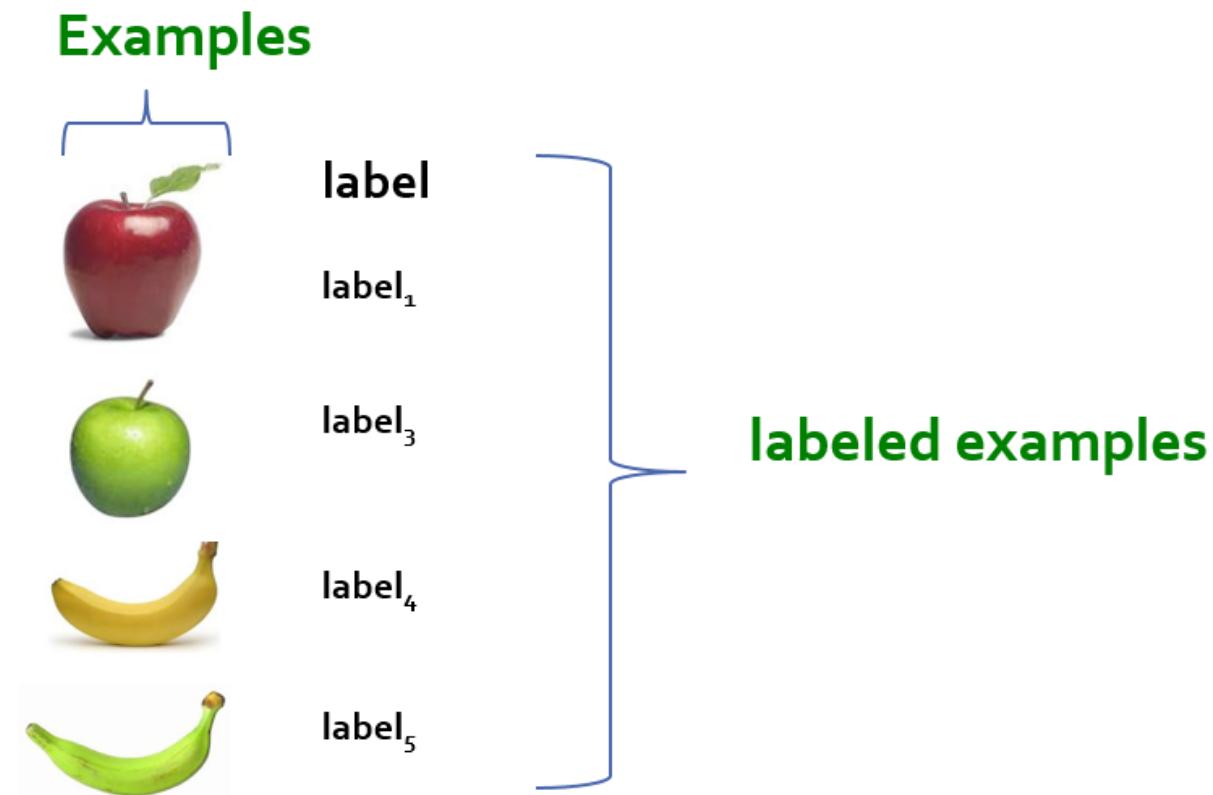
- Start training by providing a list of all the labels (or classes, or categories) that we want it to learn.
- Make the list so that it has all the labels for all the samples in the training set, with the duplicates removed.
- Train the system with lots of photos and their labels, until it does a good job of predicting the correct label for each photo.
- Now, turn the system loose on new photos it hasn't seen before.
- For those objects it **saw during training**, It should **properly label images**.
- Caution: For those objects it **did not see during training**, the system will try to pick the best category from those it knows about.
- Next Figure shows the idea.

# SL – Classification Example

---

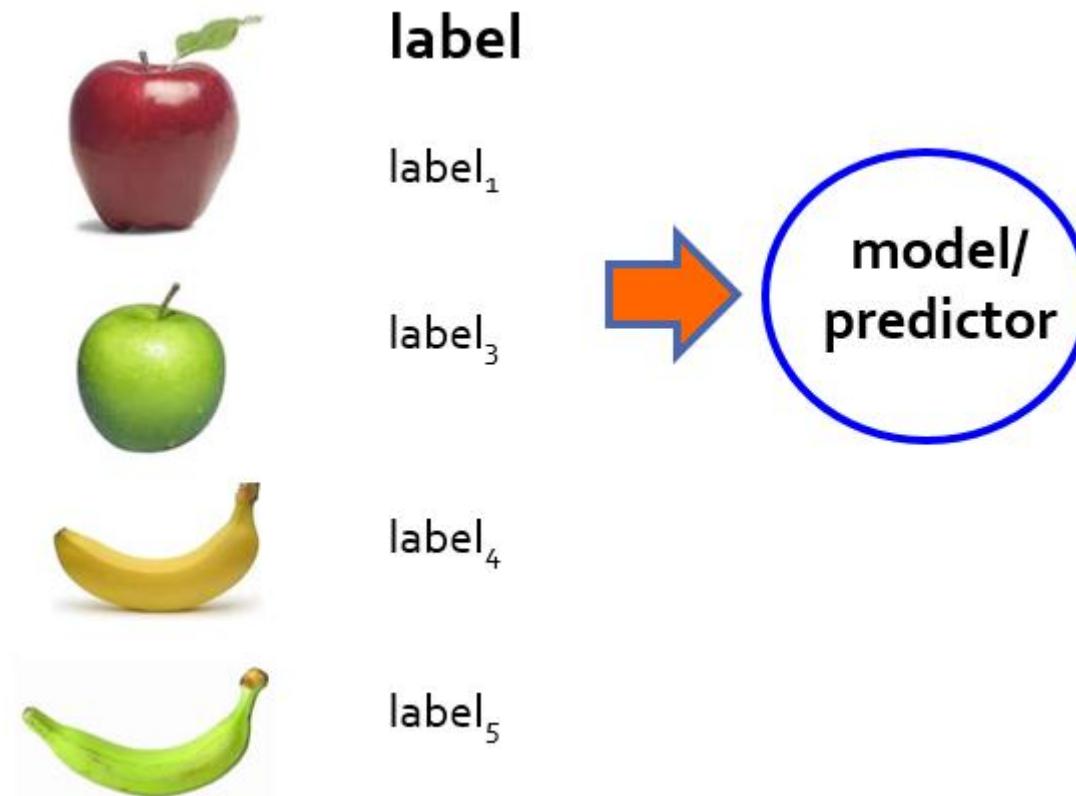
- **Example:** Sort and label photos of everyday objects.
- We want to sort them : an apple peeler, a salamander, a piano, and so on.
- We want to **classify** or **categorize** these photos.
- The process is called **classification or categorization**.

# SL – Classification Example



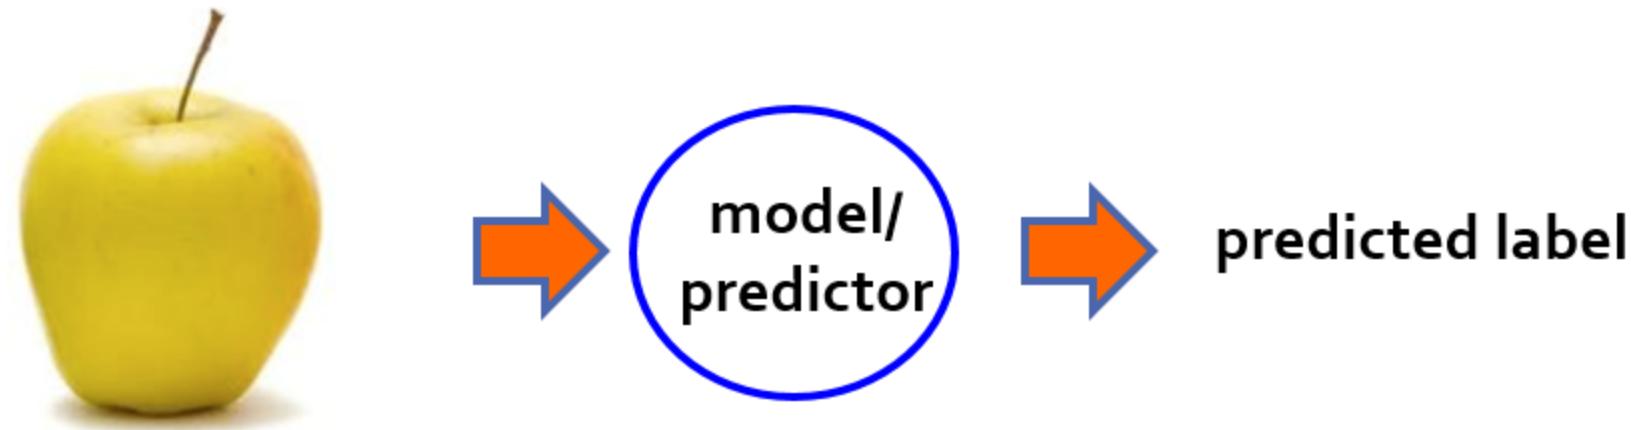
**Supervised learning: given labeled examples**

# SL – Classification Example



**Supervised learning: given labeled examples**

# SL – Classification Example



Supervised learning: learn to predict new example

# SL – Classification Example



label

apple



apple



banana

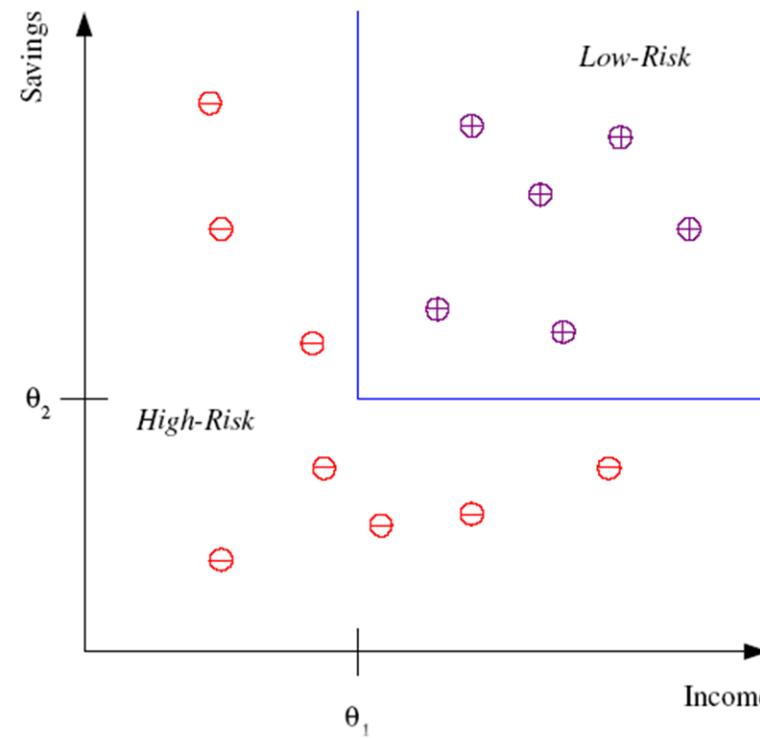


banana

**Classification: a finite set of labels**

# SL – Classification Example

Differentiate between **low-risk** and **high-risk** customers from their *income* and *savings*



# SL – Classification Example

- In Figure, we used a trained classifier to identify four images never seen before.
- The system had not been trained on metal spoons or headphones,
- In both cases it found best match it could.
- To correctly identify those objects, the system needs to see multiple examples of them during training.



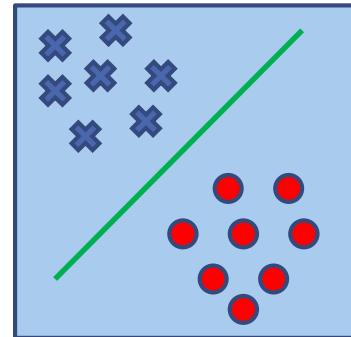
# Learning techniques

---

- Supervised learning categories and techniques
  - Linear classifier (numerical functions)
  - Parametric (Probabilistic functions)
    - Naïve Bayes, Gaussian discriminant analysis (GDA), Hidden Markov models (HMM), Probabilistic graphical models
  - Non-parametric (Instance-based functions)
    - K-nearest neighbors, Kernel regression, Kernel density estimation, Local regression
  - Non-metric (Symbolic functions)
    - Classification and regression tree (CART), decision tree
  - Aggregation
    - Bagging (bootstrap + aggregation), Adaboost, Random forest

# Learning techniques

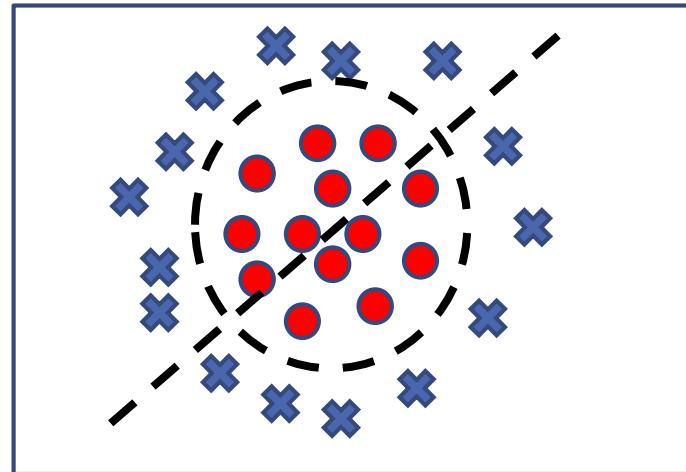
- Linear classifier



- Techniques:
  - Perceptron
  - Logistic regression
  - Support vector machine (SVM)
  - Ada-line
  - Multi-layer perceptron (MLP)

# Learning techniques

- Non-linear case



- Support vector machine (SVM):
  - Linear to nonlinear: Feature transform and kernel function

# Classification Applications

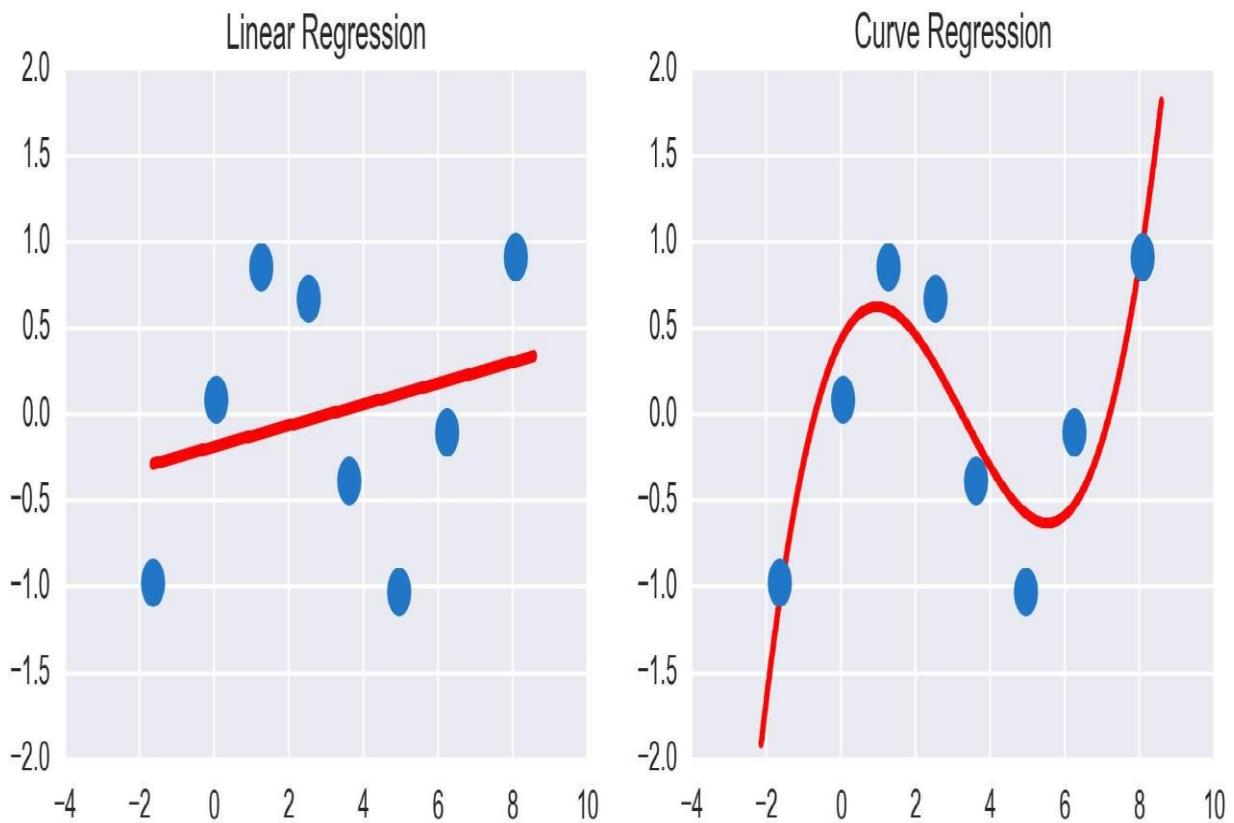
- **Face recognition:** Pose, lighting, occlusion (glasses, beard), make-up, hair style



- **Character recognition:** Different handwriting styles.
- **Speech recognition:** Temporal dependency.
- **Medical diagnosis:** From symptoms to illnesses
- **Web Advertising:** Predict if a user clicks on an ad on the Internet.
- **Biometrics:** Recognition/authentication using physical and/or behavioral characteristics: Face, iris, signature, etc

# SL - Regression

- Regression is process of filling in or predicting data .
- “Regression” uses statistical properties of the data to estimate missing or future values.
- Most famous kind of regression is **linear regression**.
- **Left:** Linear regression fits a straight line (red) to the data (blue). The line is not a very good match to the data, but it has the benefit of being simple.
- **Right:** Nonlinear regression fits a curve to same data. This is a better match to the data, but has more complicated form and requires more work (and thus more time)



# Supervised learning: Regression

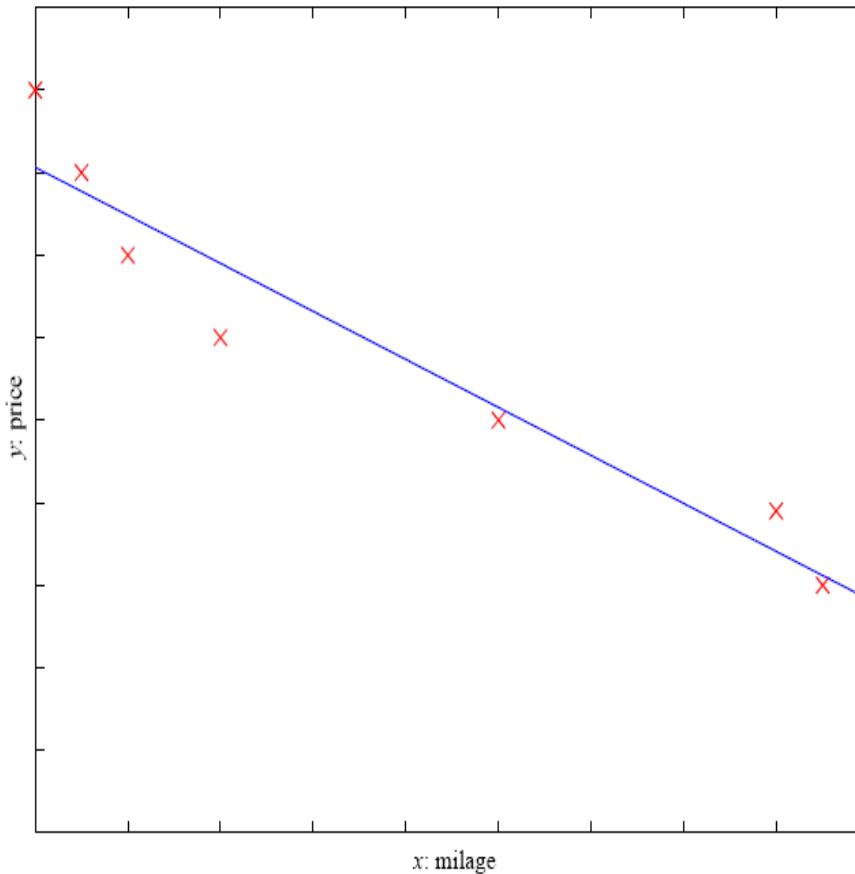
	label
	-4.5
	10.1
	3.2
	4.3

**Regression: label is real-valued**

# Regression Example

- Price of a used car

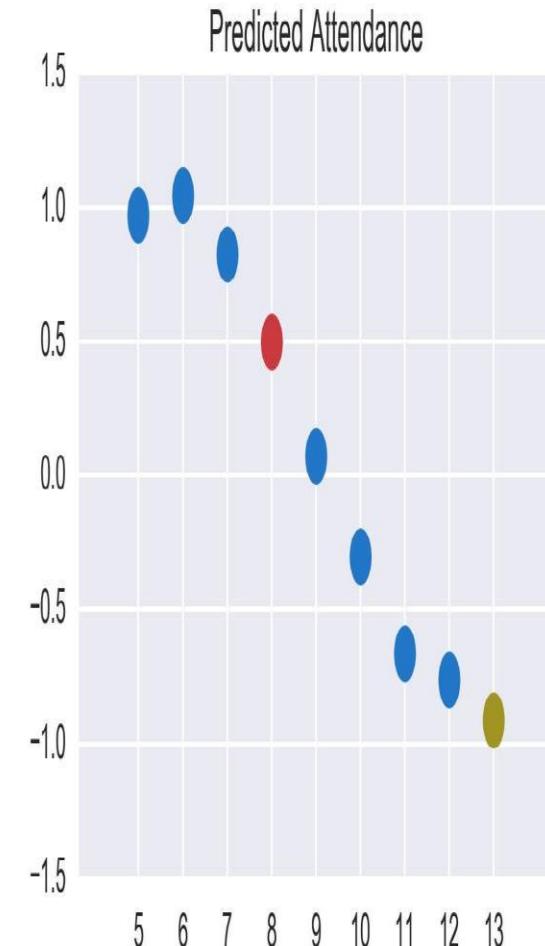
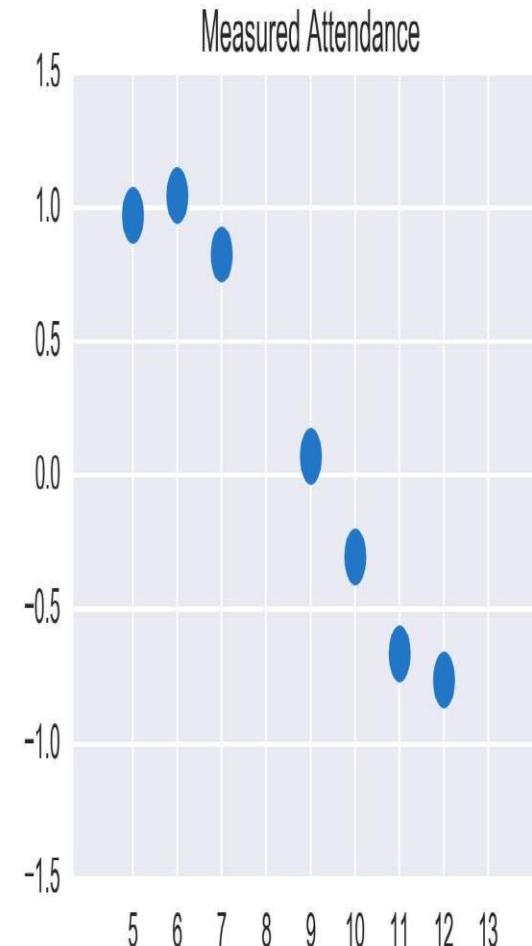
$x$  : car attributes (e.g. mileage)  
 $y$  : price



# Regression Example

## Example: Music band attendance

- Problem: An incomplete collection of measurements of attendance.
- Estimate missing values of attendance.
- Data: Attendance at a series of concerts at a local arena.
- Problem: Unfortunately, we lost count for one evening's performance.
- Also want to know what tomorrow's attendance is likely to be.



# Regression Applications

---

- Economics/Finance: predict the value of a stock
- Epidemiology
- Car/plane navigation: angle of the steering wheel, acceleration
- Temporal trends: weather over time

# Supervised learning: Ranking

	label
	1
	4
	2
	3

**Ranking: label is a ranking**

# Ranking example

- Given a query and a set of web pages, rank them according to relevance

The screenshot shows a Google search results page with the following details:

- Search Query:** machine learning
- Number of Results:** About 130,000,000 results (0.26 seconds)
- Top Result (Rank 1):**
  - Title:** Machine learning - Wikipedia, the free encyclopedia
  - URL:** en.wikipedia.org/wiki/Machine\_learning
  - Description:** Machine learning, a branch of artificial intelligence, concerns the construction and study of systems that can learn from data. For example, a machine learning ...
  - Additional Info:** Artificial intelligence - Supervised learning - List of machine learning ... - Weka
  - Interaction:** Franck Demoncourt +1'd this
- Second Result (Rank 2):**
  - Title:** CS 229: Machine Learning
  - URL:** cs229.stanford.edu/
  - Description:** Check out this year's awesome projects at Fall 2012 Projects. Come check out the cool new projects during the CS229 Poster Session this Thursday December ...
  - Interaction:** You've visited this page 2 times. Last visit: 8/14/13
- Third Result (Rank 3):**
  - Title:** Machine Learning | Coursera
  - URL:** https://www.coursera.org/course/ml
  - Description:** Machine learning is the science of getting computers to act without being explicitly programmed. In the past decade, machine learning has given us self-driving ...
  - Interaction:** Franck Demoncourt and 3 other people +1'd this
- Fourth Result (Rank 4):**
  - Title:** Machine Learning Department - Carnegie Mellon University
  - URL:** www.ml.cmu.edu/
  - Description:** Large group with projects in robot learning, data mining for manufacturing and in multimedia databases, causal inference, and disclosure limitation.
- Fifth Result (Rank 5):**
  - Title:** Machine Learning - MIT OpenCourseWare
  - URL:** ocw.mit.edu › Courses › Electrical Engineering and Computer Science
  - Description:** 6.867 is an introductory course on machine learning which gives an overview of many concepts, techniques, and algorithms in machine learning, beginning with ...

# Ranking Applications

---

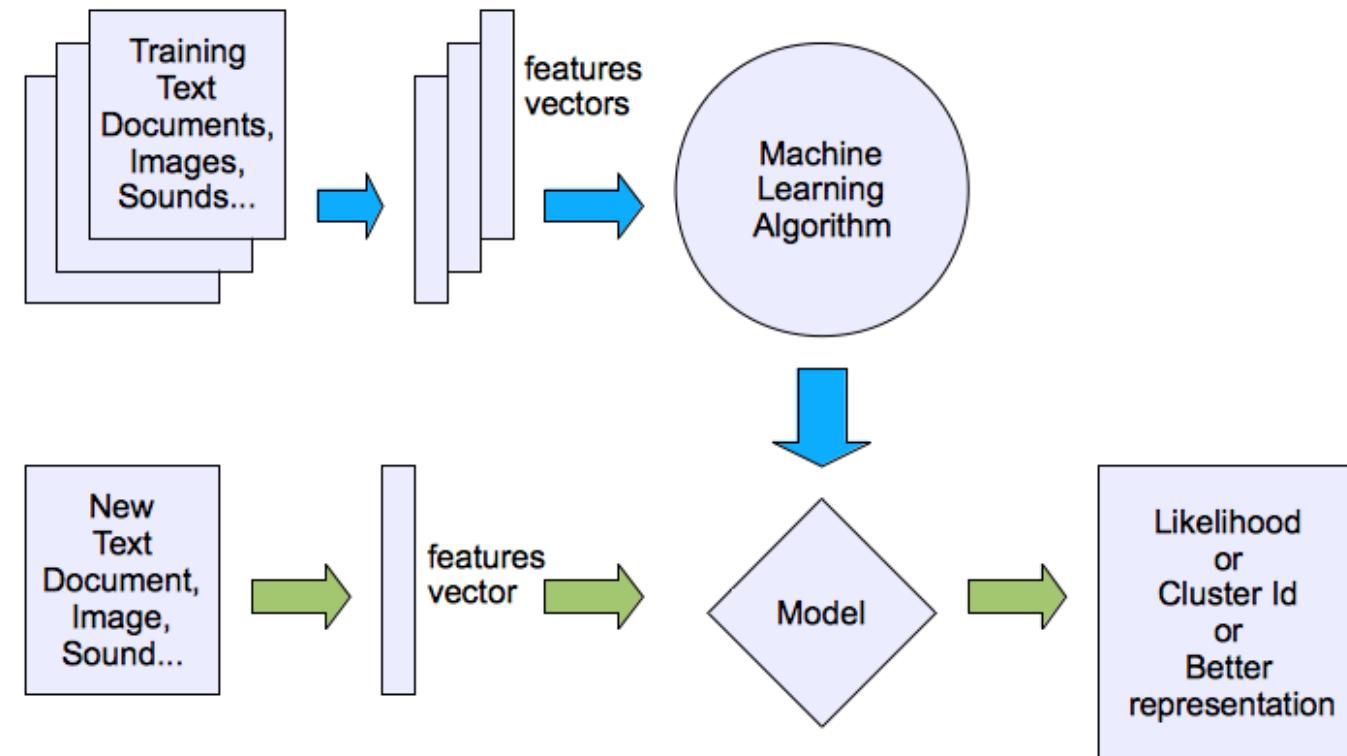
- User preference, e.g. Netflix “My List” -- movie queue ranking
- iTunes
- flight search (search in general)
- reranking N-best output lists

# Unsupervised Learning (USL) – a form of ML

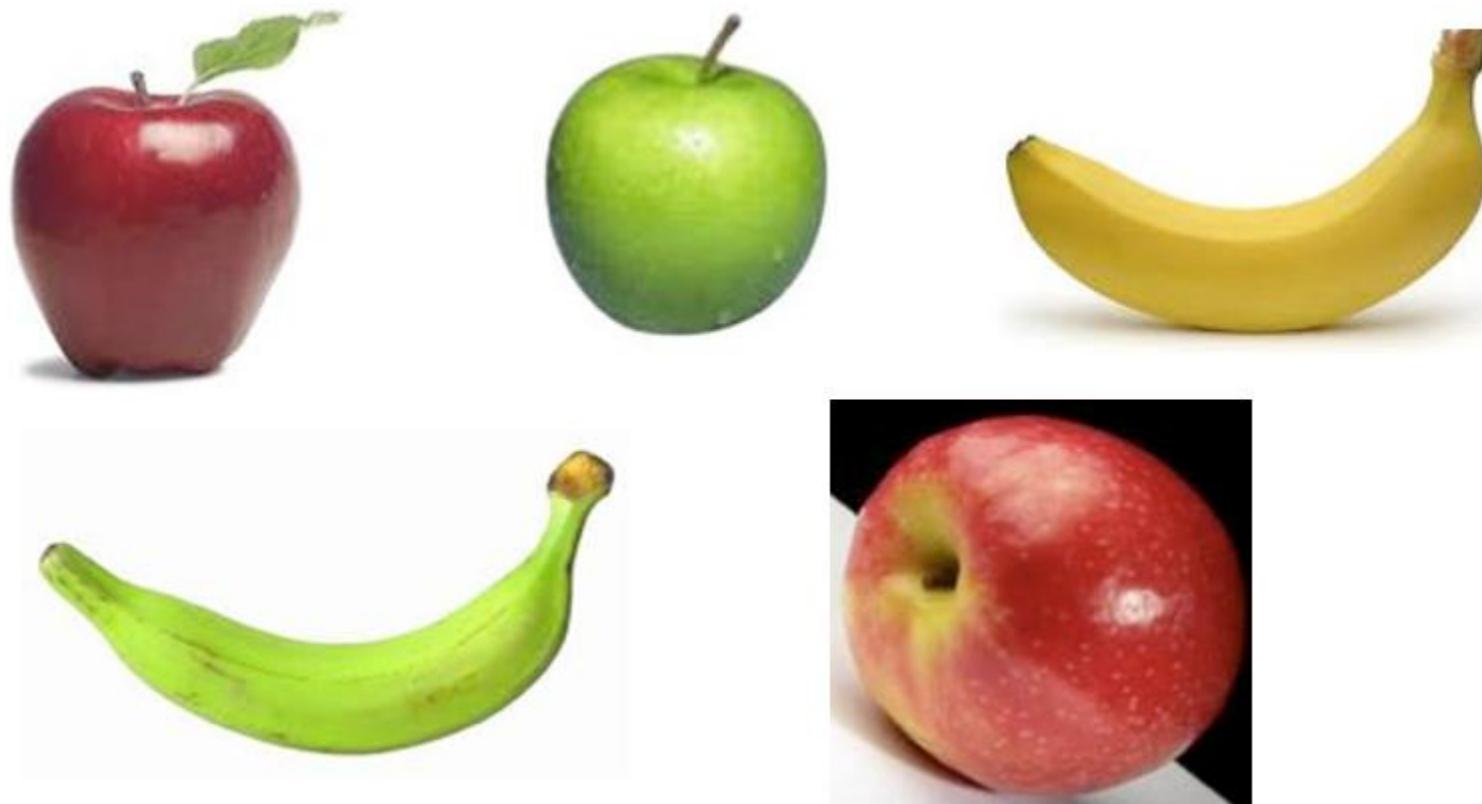
---

- When input data does not have labels, any algorithm that learns from the data belongs to USL.
- We are not “supervising” the learning process by offering labels.
- The system has to figure everything out on its own, with no help from us.
- USL is used for clustering, noise reduction, and dimension reduction.

# Unsupervised learning



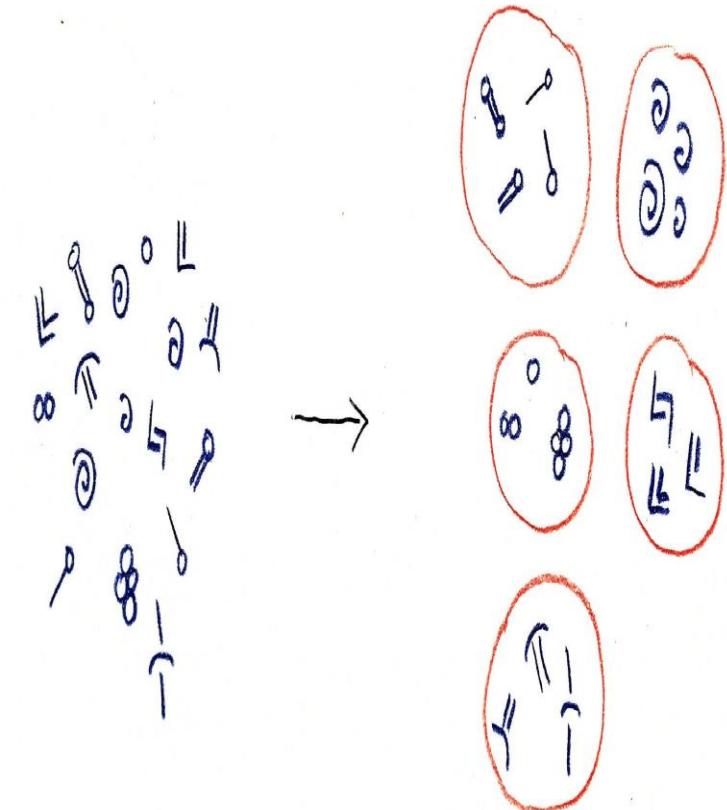
# Unsupervised learning



Unsupervised learning: given data, i.e. examples, but no labels

# USL for Clustering - Via Pottery Example

- Suppose we're digging out foundation for a new house.
- Surprise! we find the ground is filled with old clay pots and vases.
- Call an archaeologist, who says its a jumbled collection of ancient pottery, from many different places and different times.
- Archaeologist doesn't recognize any of the markings and decorations, so she can't declare for sure where each one came from.
- Some marks look like variations on same theme, while others look like different symbols.



# USL for Clustering - Via Pottery Example

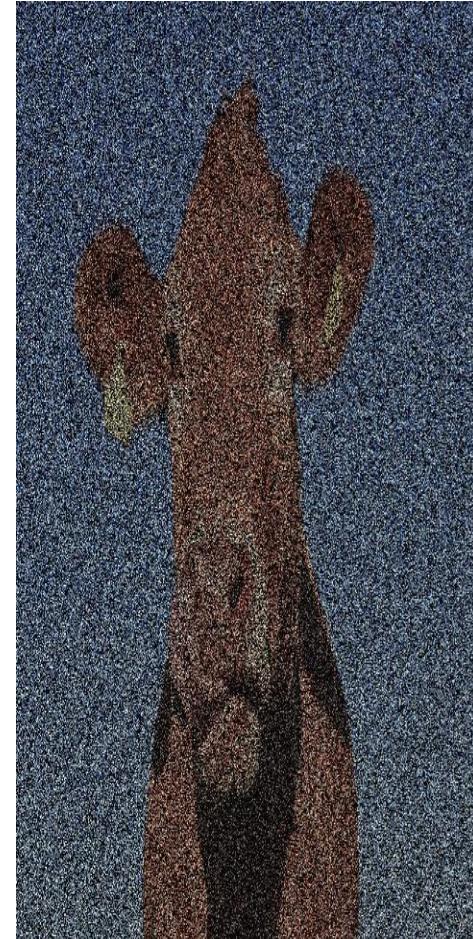
- She takes rubbings of the markings, and then tries to sort them into groups.
- But there are far too many of them for her to manage.
- She turns to a machine learning algorithm.

## Why ML?

- To automatically group the markings together in a **sensible way**.
- On the right of previous figure, we show her captured marks, and the **groupings that could be found automatically** by an algorithm.
- This is a **clustering problem**
- The ML algorithm is a clustering algorithm.
- There are many clustering algorithms to choose from.
- Because our inputs are unlabeled, this archaeologist is performing clustering, using an unsupervised learning algorithm.

# USL for Noise Reduction – Noisy Image Example

- Figure shows a noisy image, how a de-noising algorithm cleans it up.
- Why is de-noising a form of unsupervised learning (USL)?
- As we don't have labels for our data (for example, in a noisy photo we just have pixels)
- USL algorithm **estimates what part of sample is noise & removes it.**
- By removing weird and missing values from the input, learning process happens more quickly and smoothly.



# USL for Dimensionality Reduction

---

- Problem: Sometimes our samples have more features than they need.
- So, simplify data:
- Remove uninformative features,
- Combine redundant features, or
- For these tasks, there are USL algorithms that can do the job.
- USL finds a way to reduce the number of features of our data - called as **dimension reduction**.

# USL for Dimensionality Reduction Example#1

## Weather

- **Data:** Weather samples in the **desert at the height of summer.**
- Record daily wind speed, wind direction, and rainfall.
- Given the season and locale, **the rainfall value will be 0 in every sample.**
- If we use these samples in a machine learning system, the computer will need to process and interpret this useless, constant piece of information with every sample.
- At best this would **slow down the analysis.**
- At worst it could **affect the system's accuracy,** because the computer would devote some of its finite resources of time and memory to trying to learn from this unchanging feature.

# USL for Dimensionality Reduction Example#1

## Weather

- Sometimes features contain **redundant data**.
- A health clinic might take everyone's **weight in kilograms** when they walk into the door. Then when a nurse takes them to an examination room, she measures their **weight again but this time in pounds**.
- Same information repeated twice, but it might be hard to recognize that because the values are different.
- Like the useless rainfall measurements, this redundancy will not work to our benefit

# Unsupervised Learning techniques

---

- Unsupervised learning categories and techniques
  - **Clustering**
    - K-means clustering
    - Spectral clustering
  - **Density Estimation**
    - Gaussian mixture model (GMM)
    - Graphical models
  - **Dimensionality reduction**
    - Principal component analysis (PCA)
    - Factor analysis

# Unsupervised learning applications

---

- learn clusters/groups without any label
- customer segmentation (i.e. grouping)
- image compression
- bioinformatics: learn motifs
- Face detection
- Object detection and recognition
- Image segmentation
- Multimedia event detection
- Economical and commercial usage

# Semi-Supervised Learning (Generators)

---

- This process of **data generation** is implemented by ML algorithms called **generators**.
- Train generators with large numbers of examples- so that they can produce new versions with lots of variation.
- We don't need labels to train generators, so its unsupervised learning techniques.
- But we do give generators some feedback as they're learning, so they know if they're making good enough fakes for us or not.
- A generator is in the middle ground. It doesn't have labels, but it is getting some feedback from us. We call this middle ground **semi-supervised learning**.

# Semi-Supervised Learning (Generators)

## Example - Persian Carpets in Movie

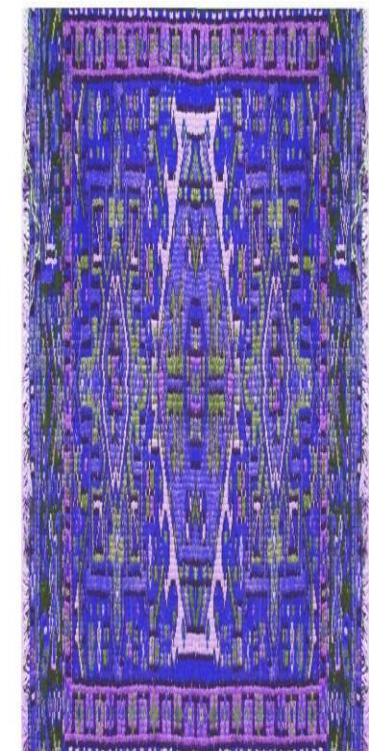
- Problem: Want hundreds of carpets, all over the warehouse.
- Carpets on the floors, carpets on the walls, and carpets in great racks in the middle of the space.
- Each carpet to look real, but be different from all the others.
- Our budget is nowhere near big enough to buy, or even borrow, hundreds of unique carpets.
- So instead, we buy just a few carpets, and then we give them to our props department to make many fake carpets.
- Figure shows a **Persian carpet** that we'd like to generalize.



# Semi-Supervised Learning (Generators)

## Example - Persian Carpets in Movie

- Figure shows a **Persian carpet** that we'd like to generalize.
- Figure shows some new fake carpets based on the starting image of previous figure made by ML Algorithms.



# Reinforcement Learning

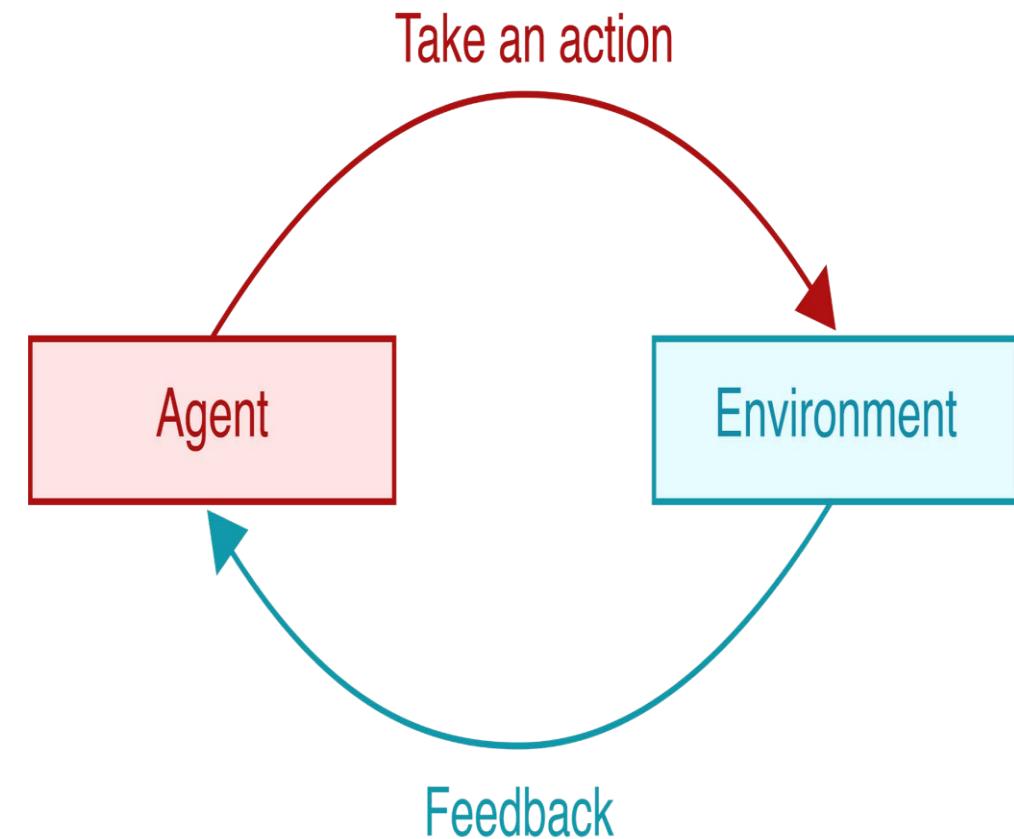
---

- Agent makes decisions and takes actions (the chef).
- Environment is everything else in the universe (the child).
- Environment gives feedback or a reward signal to agent after every action.
- Feedback tells how good or bad that action is.
- The reward signal is often just a single number, where larger positive numbers mean the action was considered better, while more negative numbers can be seen as punishments.
- The reward signal is not a label, nor a pointer to a specific kind of “correct answer.”

# Reinforcement Learning (Contd.)

In reinforcement learning,

- An agent (who acts)
- An environment (everything except agent).
- The agent acts, and the environment responds by sending feedback in the form of a reward signal.



# How's RL different from SL?

---

- The general plan of learning from mistakes is the same, but the mechanism is different.
- Supervised learning: system produces a result (typically a category or a predicted value), and then we compare it to the correct result, which we provide.
- Reinforcement learning: There is no correct result. The data has no label.
- There's just feedback that tells us how well we're doing.
- Feedback tells that our action was "good" or "bad."
- In contrast to supervised learning algorithms, the reward signal is not a label and no idea/pointer to "correct answer."

# Reinforcement learning-Example

left, right, straight, left, left, left, straight GOOD

left, straight, straight, left, right, straight, straight BAD

---

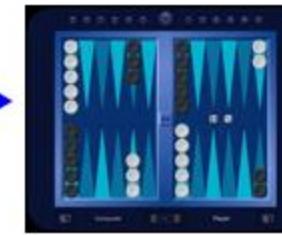
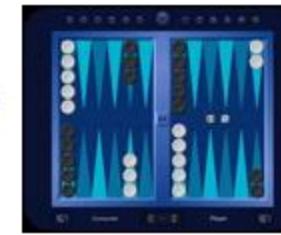
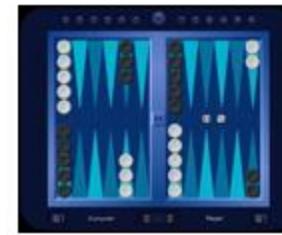
left, right, straight, left, left, left, straight 18.5

left, straight, straight, left, right, straight, straight -3

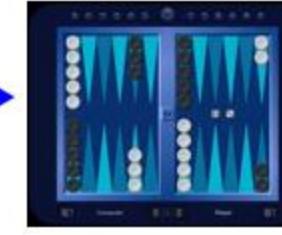
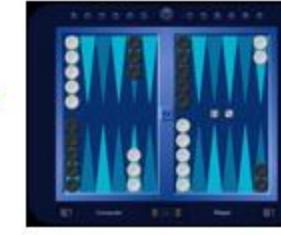
---

- Given a *sequence* of examples/states and a *reward* after completing that sequence, learn to predict the action to take in for an individual example/state

# Reinforcement learning example



**WIN!**



**LOSE!**

- Given sequences of moves and whether or not the player won at the end, learn to make good moves

# Reinforcement Learning - Example

---

- Suppose you are taking care of a friend's three-year old daughter.
- You have no idea what the young girl likes to eat.
- First dinner: Make Pasta with butter. She likes it!
- Repeat this dinner for a week. She gets bored.
- Week 2: Add some cheese, and she likes it.
- Repeat this dinner for week 2. She gets bored.
- Week 3: Try pesto sauce. But girl refuses to take a bite.
- So pasta + marinara sauce , and she rejects that too.
- Frustrated, you make a baked potato with cream. She likes it!
- Weeks 3 & 4 : Try one recipe and one variation after another, trying to develop a menu that the child will enjoy.
- Only feedback: Little girl eats the meal, or she doesn't.
- Approach to learning is **Reinforcement Learning !**

# Reinforcement Learning - Example

---

- Agent: Autonomous car
  - Environment: Traffic/people on the street.
  - Actions: Driving
  - Feedback: Driving okay if following traffic rules and keep everybody safe.
- 
- Agent: DJ at a dance club
  - Environment : Dancers
  - Feedback: Like or dislike the music.

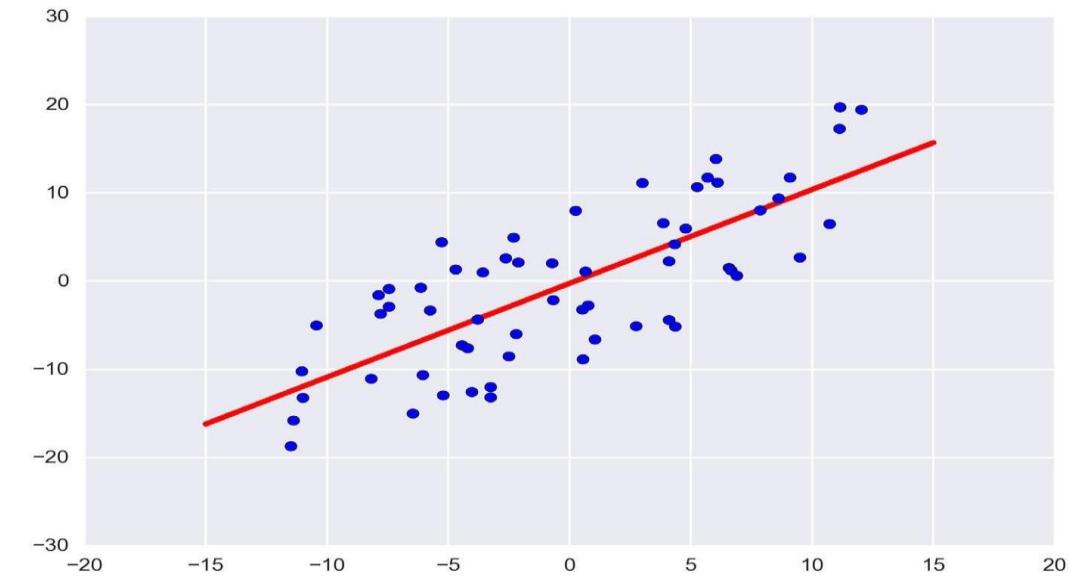
# Deep learning

---

- More recently, the phrase **deep learning** was coined to refer to
  - Approaches to machine learning that use specialized layers of computation, stacked up one after the next.
  - This makes a “deep” structure, like a stack of pancakes.
  - Deep learning (DL) refers to **the nature of the system we create, not to any particular algorithm.**
  - DL really refers to this **particular style or approach to machine learning (ML).**

# Compare ML & DL - Fit the best line

- **ML:** Find the best straight line through a bunch of data points, see Figure.
- **DL:** Given a set of data points (in blue), we can imagine a straightforward algorithm that computes the best straight line (in red) through those points.



# ML vs DL – Line fitting example

- ML: Represent a straight line with just a few parameters.
- Use some formula to compute parameter values, given the input data points.
- This is a familiar algorithm, uses analysis to find the best way to solve a problem.
- Strategy used by many ML algorithms.
- DL: Don't know how to directly calculate the right answer, so we build a system that can figure out how to do that itself.
  - For DL, we create an algorithm that can work out its own answers, rather than implementing a known process that directly yields an answer
  - DL learns slowly. Every time program sees a new piece of data, it improves its own parameters
  - DL ultimately finds a set of good values.
  - Much more open-ended than the one that fits a straight line.

# Heart of DL success

---

- What's a main reason for the enormous success of deep learning algorithms?
- DL success is due to the programs that find their own answers (apart from availability of high performance computing power).

# Evaluation Metrics

**Table: A sample test set with model predictions.**

ID	Target	Pred.	Outcome	ID	Target	Pred.	Outcome
1	spam	ham	FN	11	ham	ham	TN
2	spam	ham	FN	12	spam	ham	FN
3	ham	ham	TN	13	ham	ham	TN
4	spam	spam	TP	14	ham	ham	TN
5	ham	ham	TN	15	ham	ham	TN
6	spam	spam	TP	16	ham	ham	TN
7	ham	ham	TN	17	ham	spam	FP
8	spam	spam	TP	18	spam	spam	TP
9	spam	spam	TP	19	ham	ham	TN
10	spam	spam	TP	20	ham	spam	FP

# Misclassification rate

**Misclassification rate** =  $\frac{\text{Number incorrect predictions}}{\text{Total Predictions}}$

$$\text{Misclassification rate} = \frac{(2 + 3)}{(6 + 9 + 2 + 3)} = 0.25$$

# Binary Prediction

---

For binary prediction problems there are 4 possible outcomes:

- True Positive (TP)
- True Negative (TN)
- False Positive (FP)
- False Negative (FN)

# The structure of a confusion matrix

		Prediction	
		positive	negative
Target	positive	$TP$	$FN$
	negative	$FP$	$TN$

# A confusion matrix for the set of predictions shown in Table

		Prediction	
		'spam'	'ham'
Target	'spam'	6	3
	'ham'	2	9

# Misclassification accuracy

$$\text{classification accuracy} = \frac{(TP + TN)}{(TP + TN + FP + FN)}$$

$$\text{classification accuracy} = \frac{(6 + 9)}{(6 + 9 + 2 + 3)} = 0.75$$

# Classification accuracy

$$\text{misclassification accuracy} = \frac{(FP + FN)}{(TP + TN + FP + FN)}$$

$$\text{misclassification accuracy} = \frac{(2 + 3)}{(6 + 9 + 2 + 3)} = 0.25$$

**Thank You**